

Workloads & Scheduling

Weight : 8

1) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes (IMPORTANT)
```

Create a new deployment named `cache-deployment` in the default namespace using a custom image `redis:7.0.13`. Ensure that the deployment has the following specifications:

- Set the replica count to 2.
- Set the strategy type `RollingUpdate`
- Configure the `MaxUnavailable` field to 30% and the `MaxSurge` field to 45%.
- Deploy the `cache-deployment` deployment and ensure that all pods are in a ready state.
- Now, Perform an image upgrade to `redis:7.2.1`.
- Examine the rolling history of the deployment, and save the Total revision count to the `total-revision.txt`.

Solution:-

Step 1: Create deployment template

```
kubectl create deployment cache-deployment --image=redis:7.0.13 --replicas=2 --dry-run=client -o yml > deployment.yml
```

Step 2: Update deployment.yml, add this under spec

```
strategy:
```

```
  type: RollingUpdate
```

```
  rollingUpdate:
```

```
    maxUnavailable: 30%
```

```
    maxSurge: 45%
```

Step 3: `kubectl apply -f deployment.yml`

Step 4: image upgrade to `redis:7.2.1`

```
kubectl set image deployment/cache-deployment *=redis:7.2.1
```

OR

```
kubectl set image deployment/cache-deployment redis=redis:7.2.1
```

Here, redis(is container name)=redis:7.2.1

Step 5: rolling history

```
kubectl rollout history deployment/cache-deployment
```

Step 6: save the Total revision count to the `total-revision.txt`

```
echo 2 > total-revision.txt
```

Weight : 5

2) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

Within the default namespace, there is a web application deployment named `webapp-deployment` that relies on an environment variable that can change frequently. You need to manage this environment variable using a ConfigMap. Follow these steps:

- Create a new ConfigMap named `webapp-deployment-config-map` with the key-value pair `APPLICATION=web-app`.
- Update the deployment `webapp-deployment` to utilize the newly created ConfigMap.

Solution:-

Step 1: Create ConfigMap

```
kubectl create configmap webapp-deployment-config-map  
--from-literal=APPLICATION=web-app
```

Step 2: Edit deployment

```
kubectl edit deploy webapp-deployment
```

Step 3: Replace From-

```
spec:  
  containers:  
  - name: webapp-container  
    image: nginx:latest  
    env:  
    - name: APPLICATION  
      value: web-app
```

To -

```
spec:  
  containers:  
  - name: webapp-container  
    image: nginx:latest  
    env:  
    - name: APPLICATION  
      valueFrom:  
        configMapKeyRef:  
          name: webapp-deployment-config-map  
          key: APPLICATION
```

Weight : 2

3) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

There is `redis-deploy` deployment in `redis-ns` namespace with replica `1` , your task is to scale it to `3` .

Solution:-

```
kubectl scale deploy redis-deploy -n redis-ns --replicas=3
```

Weight : 5

4) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

Currently, the `webapp-deployment` is running with below sensitive database environment variables directly embedded in the deployment YAML. To enhance security and protect the sensitive data, perform the following steps:

- Create a Kubernetes Secret named `db-secret` with the sensitive database environment variable values(base64):
 - Key: `DB_Host` , Value: `mysql-host`
 - Key: `DB_User` , Value: `root`
 - Key: `DB_Password` , Value: `dbpassword`
- Update the `webapp-deployment` to load the sensitive database environment variables from the newly created `db-secret` Secret.

Solution:-

Step 1: Create secret

```
kubectl create secret generic db-secret --from-literal=DB_Host=mysql-host \
--from-literal=DB_User=root \
--from-literal=DB_Password=dbpassword
```

Step 2: Edit deployment

```
kubectl edit deploy webapp-deployment
```

Step 3: Replace From-

```
containers:
- env:
- name: DB_Host
  value: mysql-host
- name: DB_User
  value: root
- name: DB_Password
  value: dbpassword
```

To -

```
containers:
env:
- name: DB_Host
  valueFrom:
```

```

      secretKeyRef:
        name: db-secret
        key: DB_Host
    - name: DB_User
      valueFrom:
        secretKeyRef:
          name: db-secret
          key: DB_User
    - name: DB_Password
      valueFrom:
        secretKeyRef:
          name: db-secret
          key: DB_Password

```

Weight : 5

5) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

There is a issue in `my-app-deployment.yaml` manifest file, fix the issues, After fixing make sure pod are running state.

Note: Don't remove any specification

Solution:-

Step 1: First Change `image: nginx:latets` To `image: nginx:latest`

Step 2: Second Changes From-

```

resources:
  requests:
    memory: "1000Mi"
    cpu: "5.0"
  limits:
    memory: "100Mi"
    cpu: "0.5"

```

To- #give less than limit

```

resources:
  requests:
    memory: "100Mi"
    cpu: "0.1"
  limits:
    memory: "100Mi"
    cpu: "0.5"

```

Step 3: `kubectl apply -f my-app-deployment.yaml`

Weight : 4

6) For this question, please set this context (In exam, diff cluster name)

```
kubect1 config use-context kubernetes-admin@kubernetes
```

The deployment named `video-app` has experienced multiple rolling updates and rollbacks. Your task is to total revision of this deployment and record the image name used in 3rd revision to file `app-file.txt` in this format

```
REVISION_TOTAL_COUNT, IMAGE_NAME .
```

Solution:-

Step 1: Check the Total revision

```
kubect1 rollout history deployment video-app
```

Step 2: Check the 3rd revision image name

```
kubect1 rollout history deployment video-app --revision=3
```

Step 3: Save it in file

```
echo "3,redis:7.0.13" > app-file.txt
```

Weight : 2

7) For this question, please set this context (In exam, diff cluster name)

```
kubect1 config use-context kubernetes-admin@kubernetes
```

Fresher deployed a pod named `my-pod` . However, while specifying the resource limits, they mistakenly given `100Mi` storage limit instead of `50Mi`

- node doesn't have sufficient resources, So change it to `50Mi` only.

Solution:-

Step 1: get the pod template

```
kubect1 get pod my-pod -o yaml > pod.yaml
```

Step 2: Update pod template, From-

```
resources:
```

```
limits:
```

```
memory: 100Mi
```

```
requests:
```

```
memory: 50Mi
```

To-

```
resources:
```

```
limits:
```

```
memory: 50Mi
```

```
requests:
```

```
memory: 50Mi
```

Step 3: recreate pod

```
kubect1 replace -f pod.yaml --force
```

Weight : 6

8) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

Due to a missing feature in the current version. To resolve this issue, perform a rollback of the deployment `redis-deployment` to the previous version. After rolling back the deployment, save the image currently in use to the `rolling-back-image.txt` file, and finally increase the replica count to 3."

Solution:-

Step 1: switch to previous version

```
kubectl rollout undo deployment redis-deployment
```

Step 2: get the image name

```
kubectl describe deploy redis-deployment | grep -i image
```

Step 3: Save it in a file

```
echo "redis:7.0.13" > rolling-back-image.txt
```

Step 4: Scale up the pod

```
kubectl scale deploy redis-deployment --replicas=3
```

Weight : 2

9) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

Create a deployment named `nginx-app-deployment` using the `nginx` image and scale the application pods to 3 .

Solution:-

```
kubectl create deployment nginx-app-deployment --image=nginx --replicas=3
```

Weight : 4

10) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

You need to create a Kubernetes Pod and a Service to host a simple web application that prints "Hello, World!" when accessed. Follow these steps:

Create a Pod named app-pod with the following specifications:

- Container name: `app-container`
- Container image: `httpd:latest`

- Container port: 80

Create a Service named app-svc with the following specifications:

- Select the Pod with the label app: app-lab .
- Service port: 80
- Target port: 80
- Service type: ClusterIP
- kubectl port-forward to forward a local port to the Pod's port
- Access the web application using curl on another terminal

Solution:-

Step 1: Create pod `kubectl run app-pod --image=httpd:latest --port=80`

Step 2: Create service `kubectl expose pod app-pod --name=app-svc --port=80 --target-port=80 --type=ClusterIP`

Step 3: Port forward `kubectl port-forward svc/app-svc 8080:80`

And `curl http://localhost:8080`

Weight : 4

11) For this question, please set this context (In exam, diff cluster name)

```
kubectl config use-context kubernetes-admin@kubernetes
```

Create a pod named ubuntu-pod using the ubuntu image with the labels set to app=os . Now, create a service named ubuntu-service to expose the ubuntu-pod application within the cluster on port 8080 .

Solution:-

Step 1: Create pod `kubectl run ubuntu-pod --image=ubuntu --labels=app=os`

Step 2: Create service `kubectl expose pod ubuntu-pod --name=ubuntu-service --port=8080 --target-port=8080`
