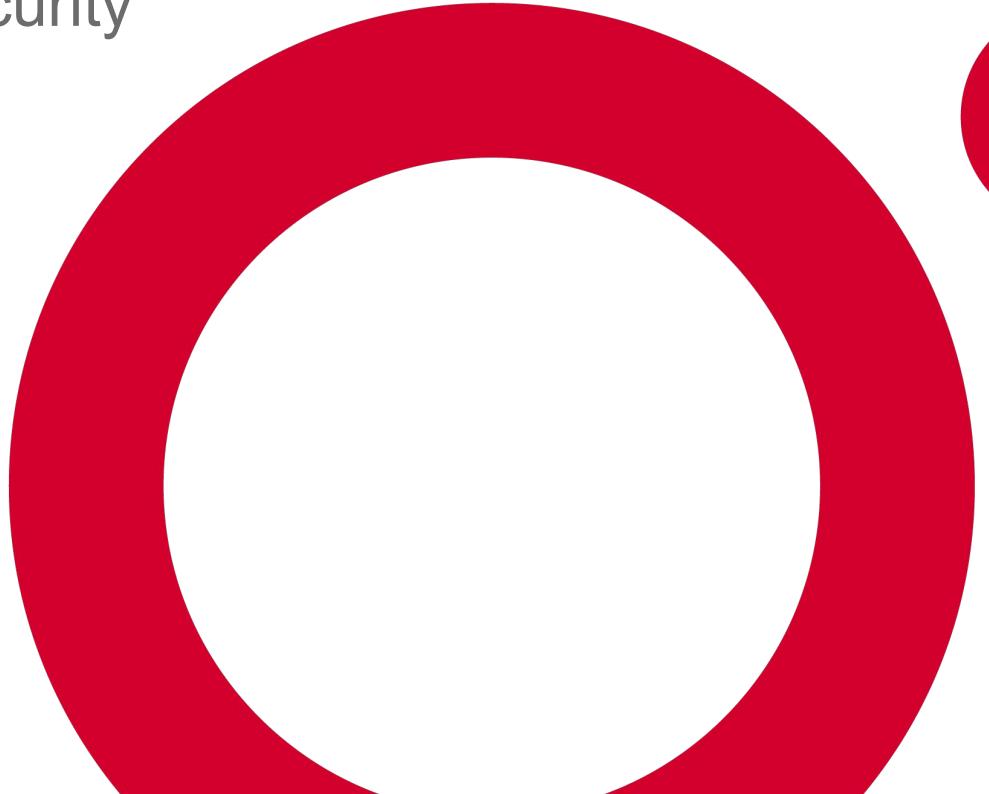




# Certified Kubernetes Security Specialist (CKS) Crash Course

*Kubernetes 1.23 Edition*



# About the trainer



**bmuschko**



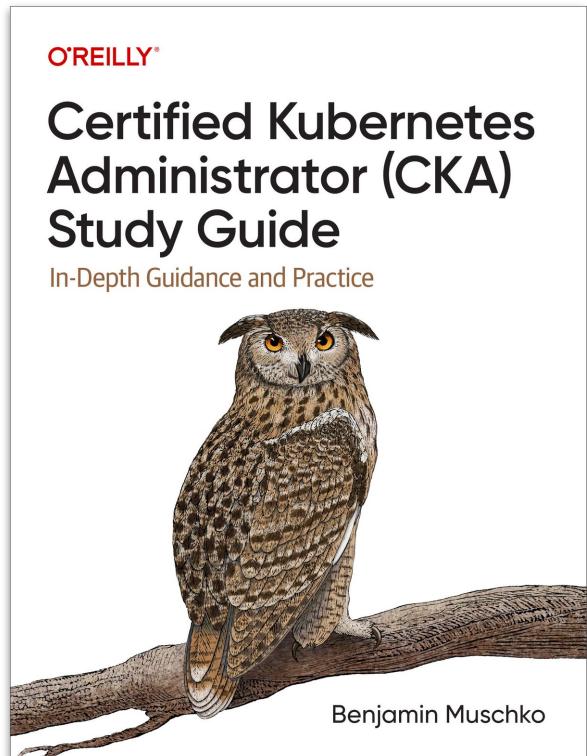
**bmuschko**



**bmuschko.com**



 **AUTOMATED  
ASCENT**  
[automatedascent.com](http://automatedascent.com)



*Companion study guide with  
practice questions*

**Released in June 2022**

Online access on O'Reilly  
learning platform:

[https://learning.oreilly.com/library/view/  
certified-kubernetes-administrator/97  
81098107215/](https://learning.oreilly.com/library/view/certified-kubernetes-administrator/9781098107215/)

# *Interactive environment*

OREILLY®

## 2-1. CKAD Pods: Creating and Interacting with a Pod

◀ Step 2 of 4 ▶

### Creating a Pod

Create a new Pod named `nginx` running the image `nginx:1.17.10`.

Expose container port 80.

The Pod should live in the namespace `ckad`.

Wait until the Pod transitions into the "Running" status.

[Solution](#)

[Continue](#)

The screenshot shows a dark-themed terminal window titled "Your Interactive Learning Environment Bash Terminal". The terminal output is as follows:

```
controlplane $ launch.sh
Waiting for kubernetes to start...
Kubernetes started
controlplane $ kubectl create namespace ckad
namespace/ckad created
controlplane $ kubectl run nginx --image=nginx:1.17.10 --restart=Never --
port=80 -n ckad
pod/nginx created
controlplane $
```

To the left of the terminal is a sidebar with the following content:

- 2-1. CKAD Pods: Creating and Interacting with a Pod**
- Creating a Pod**
- Create a new Pod named `nginx` running the image `nginx:1.17.10`.
- Expose container port 80.
- The Pod should live in the namespace `ckad`.
- Wait until the Pod transitions into the "Running" status.
- [Solution](#)
- [Continue](#)



## Kubernetes sandbox

Online access on O'Reilly learning platform:

<https://learning.oreilly.com/scenarios/kubernetes-sandbox/9781492062820/>

# Exam Details and Resources

Objectives, Environment, Time Management

# Exam Objectives

*“Demonstrate competence with securing container-based applications and Kubernetes platforms.”*



**The certification program allows users to demonstrate their competence in a hands-on, command-line environment.**

<https://www.cncf.io/certification/cks/>



# Exam Prerequisites

*Requires CKA certification*

**You need to be certified at the time of signing up for the CKS exam.**



<https://www.cncf.io/certification/cka/>



# Exam Domains & Weights

| Domain                                    | Weight |
|---|--------|
| Cluster Setup                             | 10%    |
| Cluster Hardening                         | 15%    |
| System Hardening                          | 15%    |
| Minimize Microservice Vulnerabilities     | 20%    |
| Supply Chain Security                     | 20%    |
| Monitoring, Logging, and Runtime Security | 20%    |



# The Curriculum

## 10% - Cluster Setup

- Use Network security policies to restrict cluster level access
- Use CIS benchmark to review the security configuration of Kubernetes components (etcd, kubelet, kubedns, kubeapi)
- Properly set up Ingress objects with security control
- Protect node metadata and endpoints
- Minimize use of, and access to, GUI elements
- Verify platform binaries before deploying

## 15% - Cluster Hardening

- Restrict access to Kubernetes API
- Use Role Based Access Controls to minimize exposure
- Exercise caution in using service accounts e.g. disable defaults, minimize permissions on newly created ones
- Update Kubernetes frequently

## 15% - System Hardening

- Minimize host OS footprint (reduce attack surface)
- Minimize IAM roles
- Minimize external access to the network
- Appropriately use kernel hardening tools such as AppArmor, seccomp

## 20% - Minimize Microservice Vulnerabilities

- Setup appropriate OS level security domains e.g. using PSP, OPA, security contexts
- Manage kubernetes secrets
- Use container runtime sandboxes in multi-tenant environments (e.g. gvisor, kata containers)
- Implement pod to pod encryption by use of mTLS

## 20% - Supply Chain Security

- Minimize base image footprint
- Secure your supply chain: whitelist allowed image registries, sign and validate images
- Use static analysis of user workloads (e.g. kubernetes resources, docker files)
- Scan images for known vulnerabilities

## 20% - Monitoring, Logging and Runtime Security

- Perform behavioral analytics of syscall process and file activities at the host and container level to detect malicious activities
- Detect threats within physical infrastructure, apps, networks, data, users and workloads
- Detect all phases of attack regardless where it occurs and how it spreads
- Perform deep analytical investigation and identification of bad actors within environment
- Ensure immutability of containers at runtime
- Use Audit Logs to monitor access



# Candidate Skills



**kubernetes**

Architecture & Concepts



`kubectl`

Running Commands



container runtime

Underlying Concepts



# Exam Environment

*Online and proctored exam*



*The trinity of tooling you need to be familiar with*



# Using Kubernetes Documentation

*Kubernetes docs and subdomains (see [FAQ](#))*

- Docs: <https://kubernetes.io/docs>
- GitHub: <https://github.com/kubernetes>
- Blog: <https://kubernetes.io/blog>



---

# Using Third-Party Documentation

*Third-party tools are part of the exam (see [FAQ](#))*

- Trivy: <https://github.com/aquasecurity/trivy>
- Falco: <https://falco.org/docs>
- AppArmor:  
<https://gitlab.com/apparmor/apparmor/-/wikis/Documentation>



---

# Bookmarking Documentation

*You are allowed to use bookmarks during exam*

- Create bookmarks of documentation you find important while practicing
- Only open “allowed” pages in a single tab



# Getting Help on a Command

*Render subcommands and options with --help*

```
$ kubectl create --help
Create a resource from a file or from stdin.

JSON and YAML formats are accepted.

...
Available Commands:

...
  configmap      Create a configmap from a local file, directory or literal
  value
  deployment     Create a deployment with the specified name.
  ...

Options:
  ...
```



# Zeroing in on Command Details

*Drill into object details with the `explain` command*

```
$ kubectl explain pods.spec  
KIND:     Pod  
VERSION:  v1
```

```
RESOURCE: spec <Object>
```

```
DESCRIPTION:
```

```
...
```

```
FIELDS:
```

```
...
```

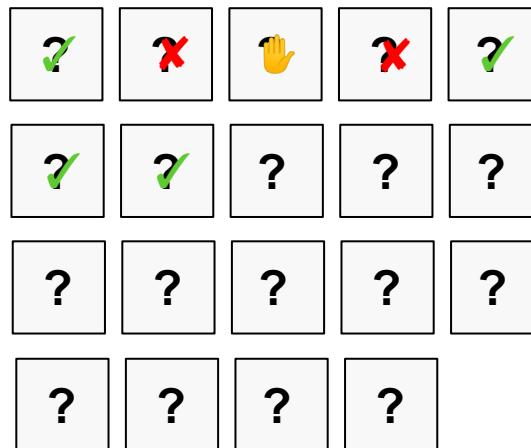
Most relevant information



---

# Time Management

*# of problems in 2 hours, use your time wisely!*



# Configuring Auto-Completion

*Allowed during exam, configurable on-demand*

```
$ kubectl cre<tab>
```

```
$ kubectl create
```



<https://kubernetes.io/docs/tasks/tools/included/optional-kubectl-configs-bash-linux/>



# Using an Alias for kubectl

*Already set up in exam environment (see [FAQ](#))*

```
$ alias k=kubectl  
$ k version  
...
```



# Setting Namespace for a Context

*Questions will ask you to run a command on a specific cluster - Make sure to execute it!*

```
$ kubectl config set-context <context-of-question>--  
  --namespace=<namespace-of-question>  
$ kubectl config use-context <context-of-question>
```



# Internalize Resource Short Names

*Some API resources provide a shortcut*

```
$ kubectl get ns
```

Usage of `ns` instead  
of namespaces

```
$ kubectl describe pvc claim
```

Usage of `pvc` instead of  
`persistentvolumeclaim`



---

# Deleting Kubernetes Objects

*Don't wait for a graceful deletion of objects...*

```
$ kubectl delete pod nginx --force
```



# Understand and Practice bash

*Practice relevant syntax and language constructs*

```
$ if [ ! -d ~/tmp ]; then mkdir -p ~/tmp; fi; while true; do echo $(date) >> ~/tmp/date.txt; sleep 5; done;
```



# Finding Object Information

*Filter configuration with context from a set of objects*

```
$ kubectl describe pods | grep -C 10 "author=John Doe"  
$ kubectl get pods -o yaml | grep -C 5 labels:
```

grep is your friend!



---

# How to Prepare

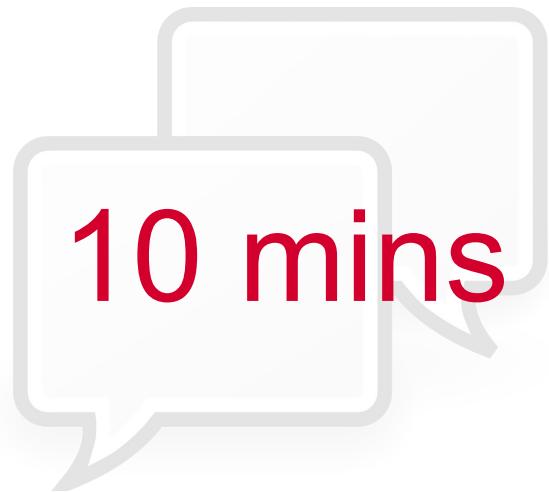
*Practice, practice, practice!*

**The key to cracking the exam**



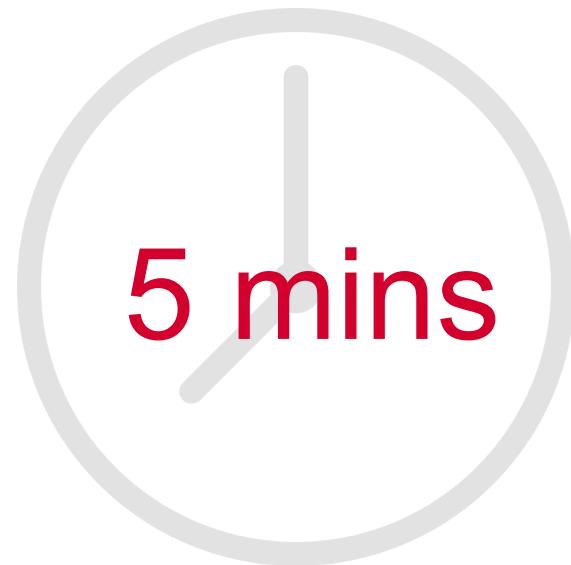
---

# Q & A



---

# BREAK

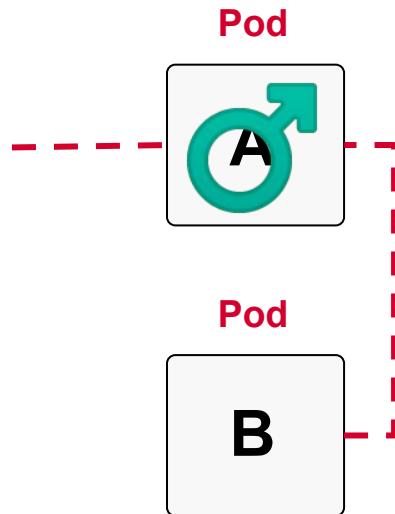


# Cluster Setup

Network Security Policies, CIS Benchmark,  
Protecting nodes and GUI, Ingress, Verifying  
Binaries

# Scenario: Gaining Access to Pod

*Once a Pod can be compromised, it can easily access other Pods*

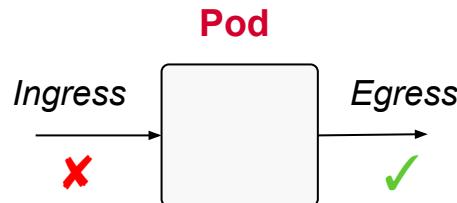


# Network Policy Rules



tier: frontend

*Which Pod does the rule apply to?* ✓



*Which direction of traffic?  
Who is allowed?*



# Creating a Network Policy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: my-network-policy
spec:
  podSelector:
    matchLabels:
      tier: frontend
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        ...
  egress:
    - to:
        ...
```



*Label selection for Pods*

*Inbound/outbound traffic*

*Who can connect to Pod?  
Where can Pod connect to?*



# General Rule of Thumb

*Start by denying all access and allowing access as needed*

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny
spec:
  podSelector: {} ← Applies to all Pods
  policyTypes:
    - Ingress ← Inbound/outbound
    - Egress   traffic is blocked
```

<https://kubernetes.io/docs/concepts/services-networking/network-policies/#default-deny-all-ingress-and-all-egress-traffic>



# Behavior of from/to Selectors

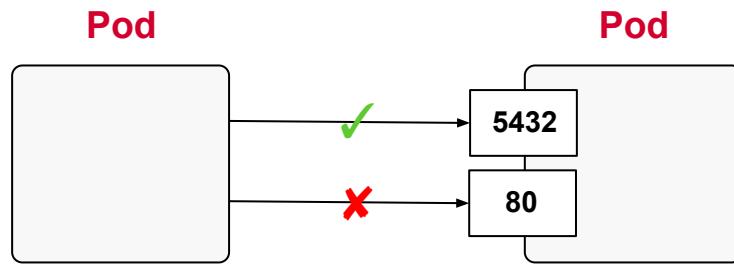
*Select by Namespace, Pod and IP address*

```
...
ingress:
- from:
  - namespaceSelector:
    matchLabels:
      app: awesome
  - podSelector:
    matchLabels:
      tier: backend
...
...
```

*Allow incoming traffic from  
Pod that matches the  
label `tier=backend` in  
the namespace `awesome`*



# Restricting Access to Ports



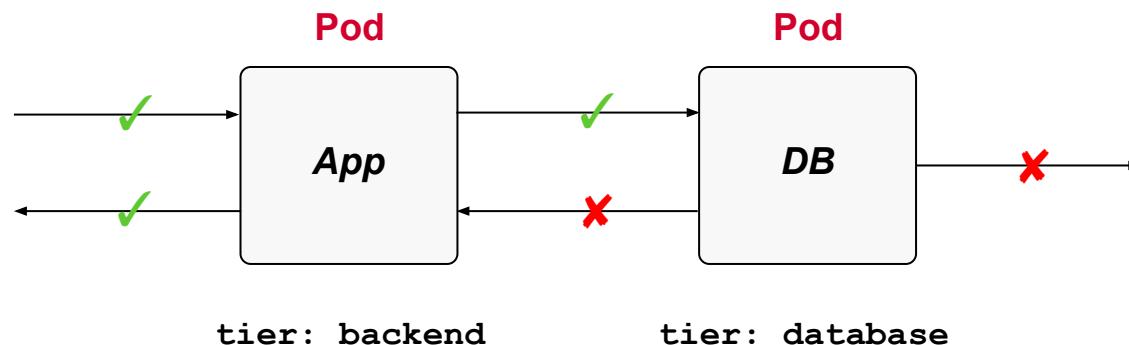
*By default all ports are open*

```
...  
ingress:  
- from:  
  - podSelector:  
    matchLabels:  
      tier: backend  
  ports:  
    - protocol: TCP  
    port: 5432  
...
```



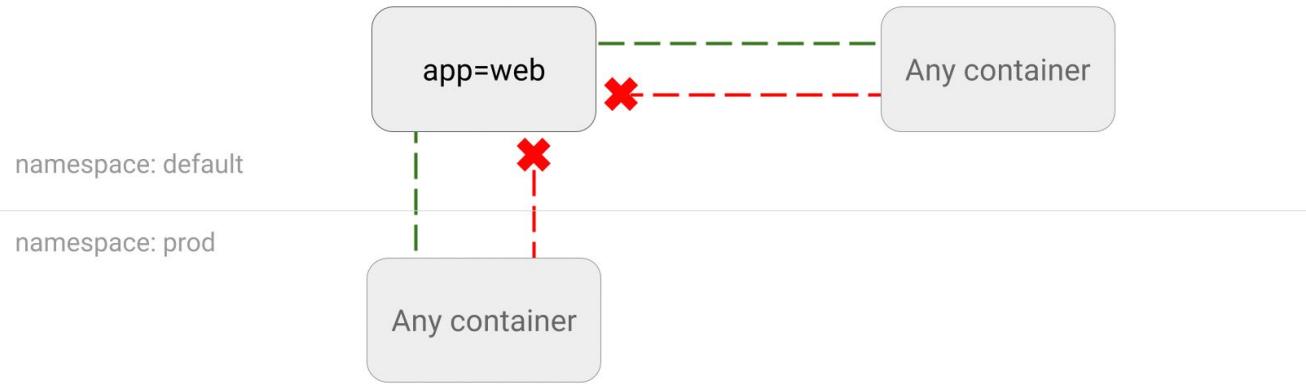
# Representative Use Case

*“Application makes request to database but database cannot make any outgoing requests.”*



# Additional Learning Resource

*Network Policies explained by use case and visualization*



<https://github.com/ahmetb/kubernetes-network-policy-recipes>



---

# EXERCISE

Restricting  
Communication  
Between Pods



# CIS Benchmarks

*Guidelines and benchmark tests for best practices in securing code*

- Published by non-profit [Center of Internet Security \(CIS\)](#)
- [CIS Kubernetes Benchmark](#) provides tests for Kubernetes environments



# Using kube-bench

Tool by Aqua for identifying and reporting issues (based on CIS)

```
$ kubectl apply -f  
https://raw.githubusercontent.com/aquasecurity/kube-benc  
h/main/job.yaml  
job.batch/kube-bench created  
  
$ kubectl get pods  
NAME             READY   STATUS    RESTARTS   AGE  
kube-bench-8f6qh   0/1     Completed   0          45s  
  
$ kubectl logs kube-bench-8f6qh  
...
```



# Results produced by kube-bench

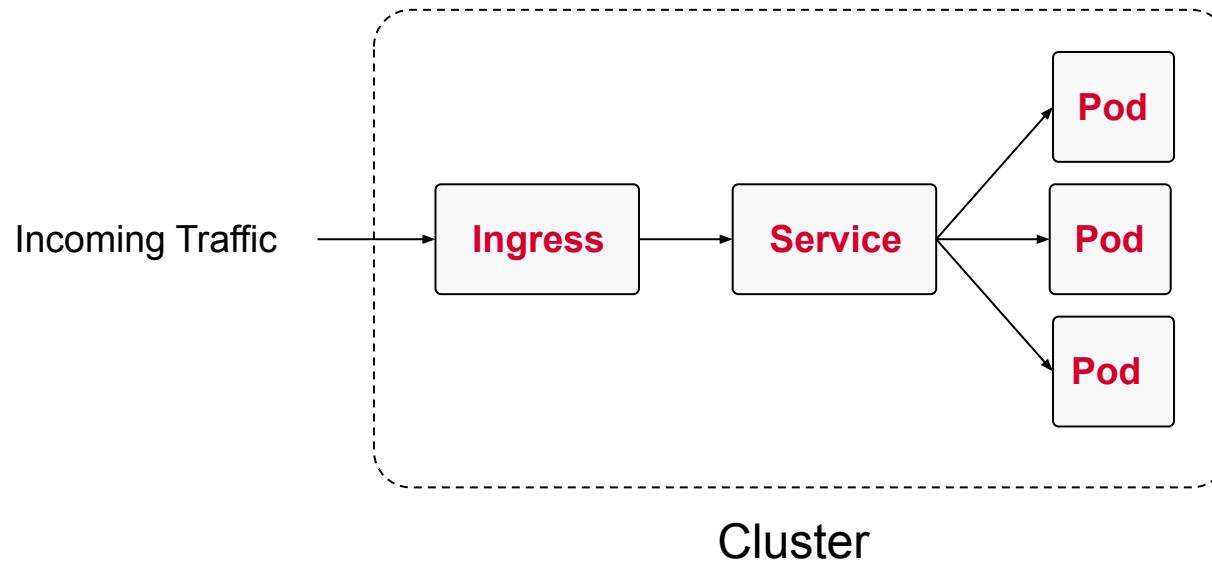
*Tests for different types of nodes and Pods running on them*

```
[INFO] 1 Master Node Security Configuration
[INFO] 1.1 Master Node Configuration Files
[PASS] 1.1.1 Ensure that the API server pod specification file permissions are
set to 644 or more restrictive (Automated)
[INFO] 1.2 API Server
[FAIL] 1.2.5 Ensure that the --kubelet-certificate-authority argument is set as
appropriate (Automated)
...
== Summary total ==
61 checks PASS
14 checks FAIL
48 checks WARN
0 checks INFO
```



# Understanding Ingress

*Manages external access to the services in a cluster via HTTP(S)*



---

# Detailed Coverage

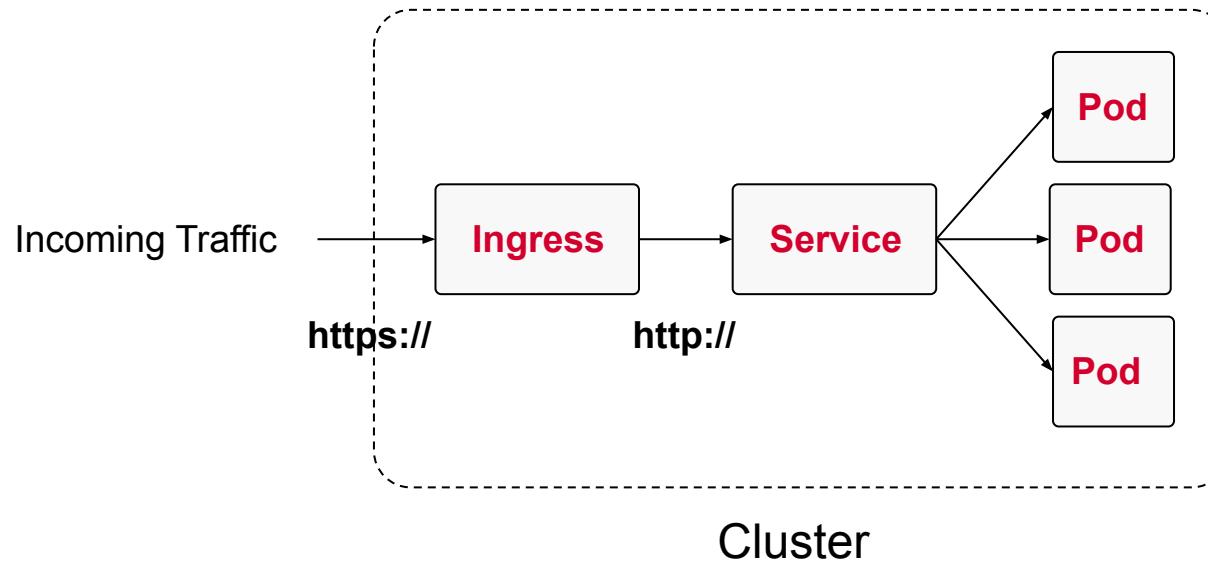
*The CKA exam lays the foundation for Ingress*

- [Certified Kubernetes Administrator \(CKA\) Study Guide](#)
- Exercise: [Creating an Ingress](#)



# Ingress with TLS termination

*Use HTTPS for Ingress, route to Service with HTTP*



# Creating the Secret

*Requires the type kubernetes.io/tls*

## Create TLS certificate and key

```
$ openssl req -nodes -new -x509 -keyout tls-ingress.key -out  
tls-ingress.crt -subj "/CN=ingress.tls"
```

## Create the Secret

```
$ kubectl create secret tls tls-ingress-secret  
--cert=tls-ingress.crt --key=tls-ingress.key
```



# Creating the Ingress

*Assign Secret upon creation*

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: tls-ingress
spec:
  tls:
  - hosts:
    - bar.foo.com
    secretName: tls-ingress-secret
  rules:
  - host: bar.foo.com
  ...
```

```
$ kubectl create ingress tls-ingress
--rule="bar.foo.com/=web-app-service
:8080,tls=tls-ingress-secret"
```



---

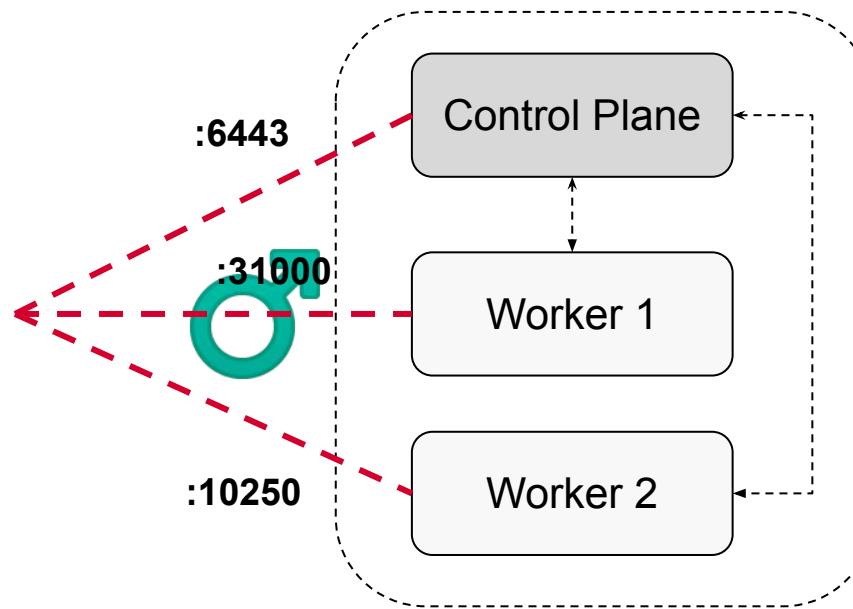
# EXERCISE

Creating an Ingress  
with TLS  
termination



# Scenario: Node Port Scans

*Nodes need to expose ports for communication*



# Default Incoming Node Ports

## Control Plane Node

| Port Range | Purpose                 |
|------------|-------------------------|
| 6643       | Kubernetes API server   |
| 2379–2380  | etcd server client API  |
| 10250      | Kubelet API             |
| 10259      | kube-scheduler          |
| 10257      | kube-controller-manager |

## Worker Node(s)

| Port Range  | Purpose           |
|-------------|-------------------|
| 10250       | Kubelet API       |
| 30000–32767 | NodePort Services |



---

# Securing Node Endpoints

*Only expose ports you need*

- Minimize attack surface area
- Limit exposure of ports
- Instantiate firewall rules to protect against access of endpoints



---

# Securing Node Metadata

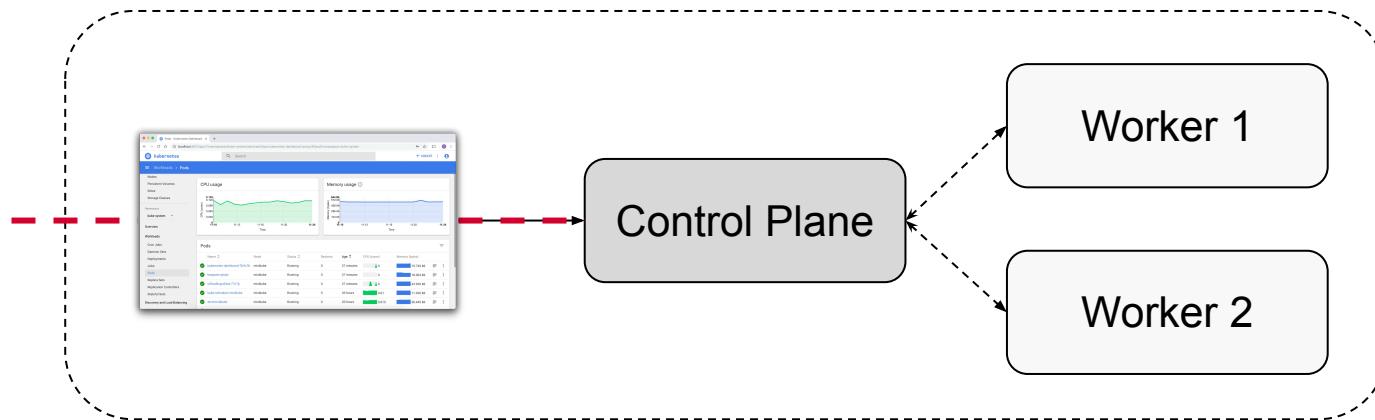
*Retrieving cluster metadata via API endpoint*

- Cloud provider Kubernetes environments can expose metadata via an API endpoint (including cloud credentials)
- Those endpoints are accessible by Pods running on the same node
- Limit permissions given to instance credentials and stand up network policies to restrict Pod access



# Scenario: Attack on GUI

*Dashboards grant access the control plane node*



---

# Securing GUI Elements

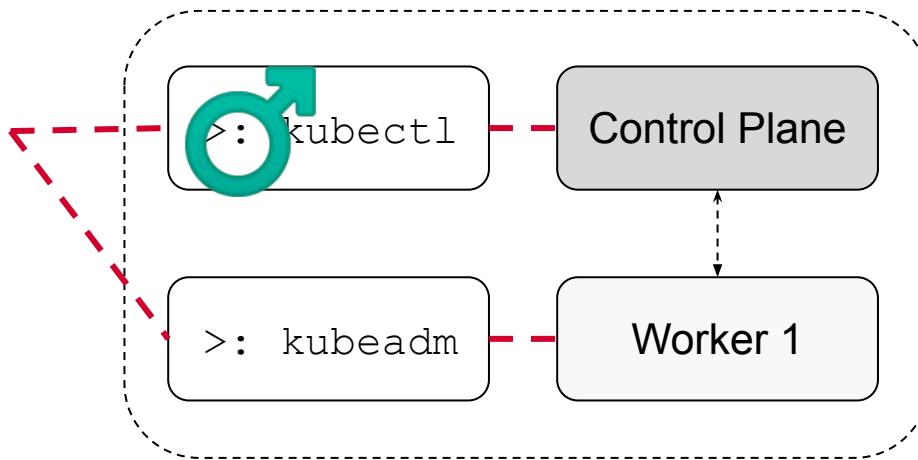
*Convenient management of Kubernetes objects*

- Exposing GUIs for Kubernetes provide attackers with a way to control the cluster
- Limit access to those GUIs (e.g. with firewall rules, RBAC permissions)
- Some organizations go as far as removing all GUI access to minimize attack surface



# Scenario: Malicious Code Inject.

*Gain control to cluster by modifying the byte code*



# Verifying Kubernetes Binaries

*Use checksum to ensure that binary hasn't been tampered with*

- Typical Kubernetes binaries are `kubectl`, `kubeadm`, and `kubelet`
- SHA256 checksum file should match the generated checksum of binary
- If both values don't match, then the file has been modified e.g. malicious code has been injected



# Retrieving the Checksum File

*Determine version, then use download link*

**Determine version of binary**

```
$ kubectl version --client --short # DO NOT RUN THIS!
Client Version: v1.23.5
```

**Download corresponding checksum file**

```
$ curl -LO
"https://dl.k8s.io/v1.23.5/bin/linux/amd64/kubectl.sha256"
```



# List of Download URLs

*Replace version, OS, and architecture to reflect your platform*

- kubectl: <https://dl.k8s.io/v1.23.5/bin/linux/amd64/kubectl.sha256>
- kubeadm: <https://dl.k8s.io/v1.23.5/bin/linux/amd64/kubeadm.sha256>
- kubelet: <https://dl.k8s.io/v1.23.5/bin/linux/amd64/kubelet.sha256>
- kube apiserver:  
<https://dl.k8s.io/v1.23.5/bin/linux/amd64/kube-apiserver.sha256>



# Verifying the Binary

*Use sha256 tool to compare binary checksum ([docs](#))*

## Linux

```
$ echo "$(cat kubectl.sha256)  kubectl" | sha256sum --check
```

## MacOSX

```
$ echo "$(cat kubectl.sha256)  kubectl" | shasum -a 256 --check
```



---

# EXERCISE

Verifying Existing  
Kubernetes  
Binaries



# Cluster Hardening

Restricting API Access, RBAC for Minimizing  
Exposure, Service Accounts, Keeping Kubernetes  
Updated

# RBAC High-Level Overview

*Three key elements for understanding concept*

## Subject

Groups  
Users  
Service-  
Accounts

## API Resources

ConfigMap  
Pod  
Deployment  
Node  
...

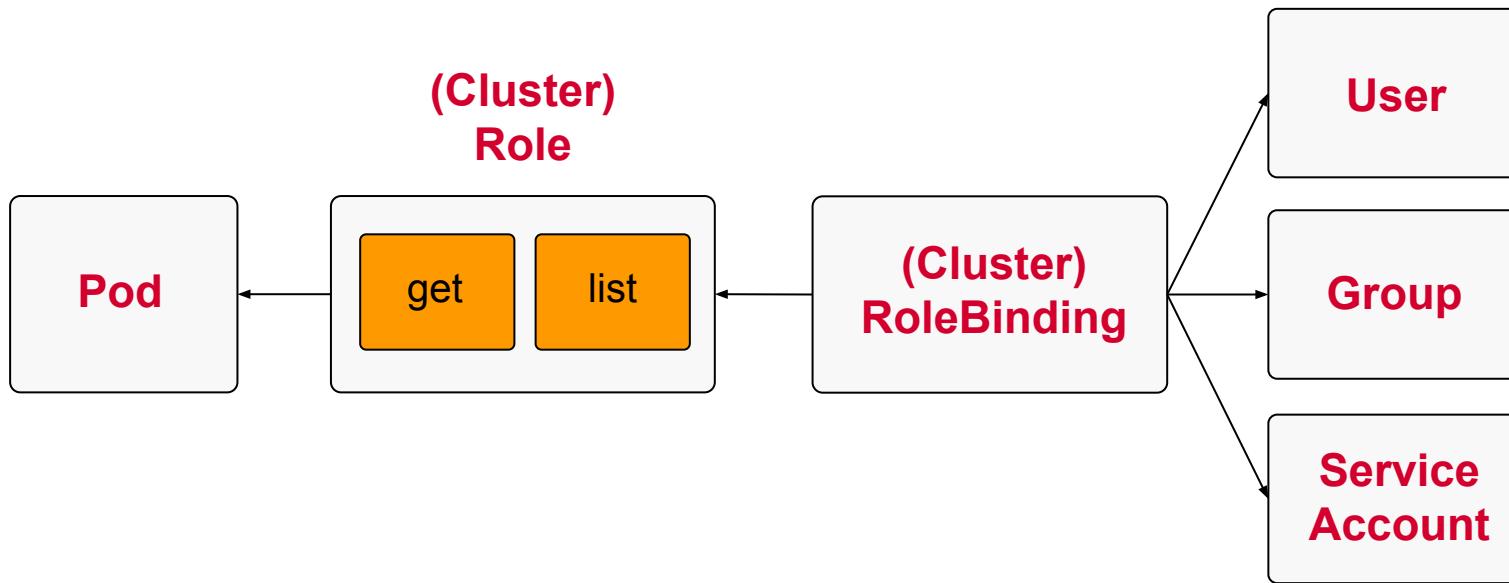
## Operations (Verbs)

create  
list  
watch  
delete  
...



# Involved RBAC Primitives

*Restrict access to API resources based on user roles*



---

# Detailed Coverage

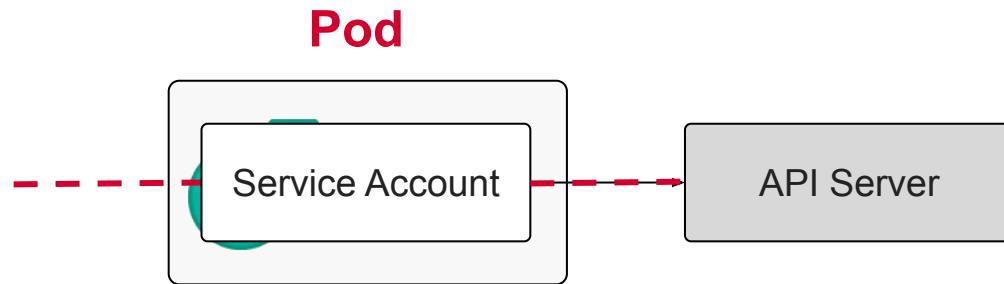
*The CKA exam lays the foundation for RBAC*

- [Certified Kubernetes Administrator \(CKA\) Study Guide](#)
- Exercise: [RBAC for a User](#)



# Scenario: Pod + Service Account

*A Service Account with too many permissions that can be misused*



# Service Account Permissions

*Restrict as much as possible*

- Service Account should only allow be allowed the permissions needed for the process running in Pod
- Restricting those permissions is implemented via (Cluster)Role and (Cluster)RoleBinding (combination of verbs + API resources)
- An attacker of a compromised Pod using the Service Account gains access to operations executed against the API server (e.g. delete Pods)



---

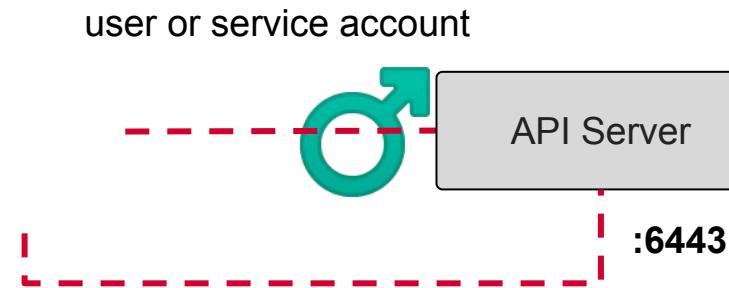
# EXERCISE

Minimizing Service  
Account Permissions  
and Usage



# Scenario: Attack on API server

*Used compromised user/service account or by port 6443*



---

# Securing API Endpoints

*Restrict as much as possible*

- Limit outside network access by defining firewall rules
- Only expose API endpoints via intranet
- User permissions should be limited and on point in case they get compromised



# Scenario: Exploiting Vulnerability

*Older Kubernetes versions can expose critical issues*



1.7

[https://www.cvedetails.com/vulnerability-list/vendor\\_id-15867/product\\_id-34016/Kubernetes-Kubernetes.html](https://www.cvedetails.com/vulnerability-list/vendor_id-15867/product_id-34016/Kubernetes-Kubernetes.html)



---

# Updating Kubernetes Frequently

*Upgrade early and often!*

- Versioning schema follows semantic versioning e.g. 1.23.5
- Check discovered vulnerabilities and upgrade to latest patch and minor version
- Make concerted effort to upgrade to next major version (they may contain breaking changes)



---

# Detailed Coverage

*The CKA exam lays the foundation for cluster management*

- [Certified Kubernetes Administrator \(CKA\) Study Guide](#)
- Exercise: [Upgrading a cluster version](#)



# System Hardening

Minimizing OS Footprint, IAM Roles, and External Network Access, Kernel Hardening Tools

# Container Namespaces

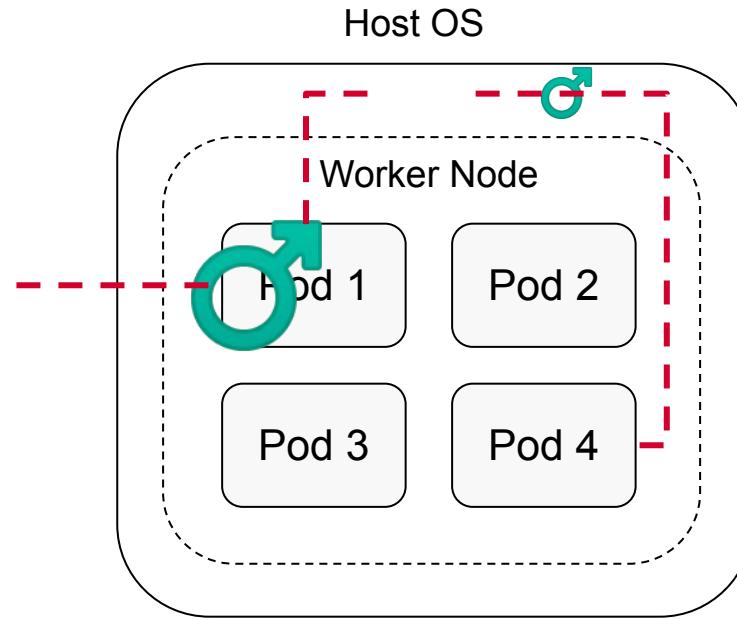
*Containers only have limited access to OS namespace*

- Regular processes run in OS namespace
- Containers run in container namespace isolated from OS namespace
- This leaves less room for attackers to gain access to host
- **Warning:** Containers can be configured to use OS namespace



# Scenario: Pod to Host OS Attack

*Compromised Pod gains access to other Pods via host OS*



# Enabling Host Namespace Usage

The Pod spec defined a handful of attributes for configuration

| YAML Element     | Purpose  |
|------------------|--|
| spec.hostIPC     | The host's inter-process communication namespace |
| spec.hostNetwork | The host's network namespace                     |
| spec.hostPID     | The host's process ID namespace                  |

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine
spec:
  hostIPC: true
  hostNetwork: true
  hostPID: true
  containers:
  - name: alpine
    image: alpine/curl:3.14
```



# Enabling Privileged Mode

*Access to host-level resources can be granted via security context*

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine
spec:
  containers:
  - name: alpine
    image: alpine/curl:3.14
    securityContext:
      privileged: true
```

Equivalent to granting root access on host. Avoid if not absolutely necessary!



---

# Granting Host Access to Pods

*Only grant those permissions if needed*

- By default, host access and privileged mode are disabled
- Enabling any of those settings opens an attack vector to the host
- Host access and privileged mode should only be enabled if really required



# IAM Roles on

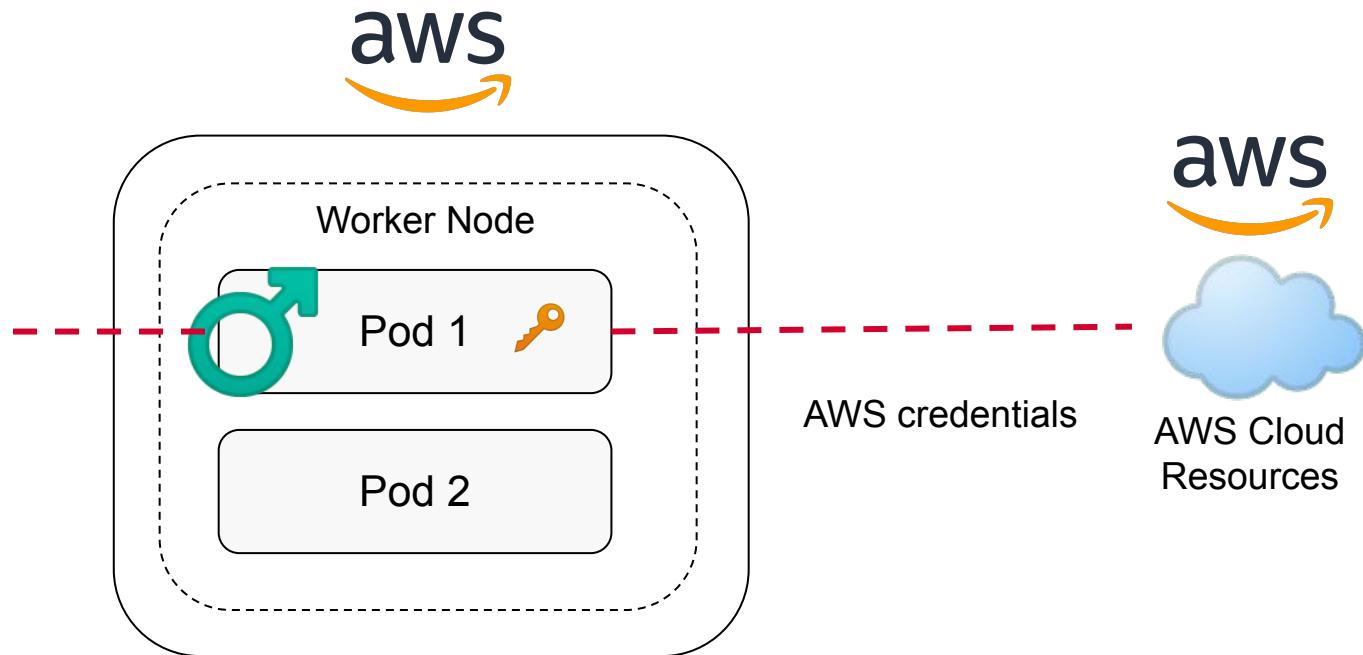
*Obtain permissions/credentials to access AWS resources*

- Only relevant if Kubernetes runs on AWS
- Containers can retrieve credentials if run on AWS and gain access to other resources



# Scenario: Misusing AWS creds.

*Compromised Pod can gain access to other AWS resources*



# Handling IAM Roles

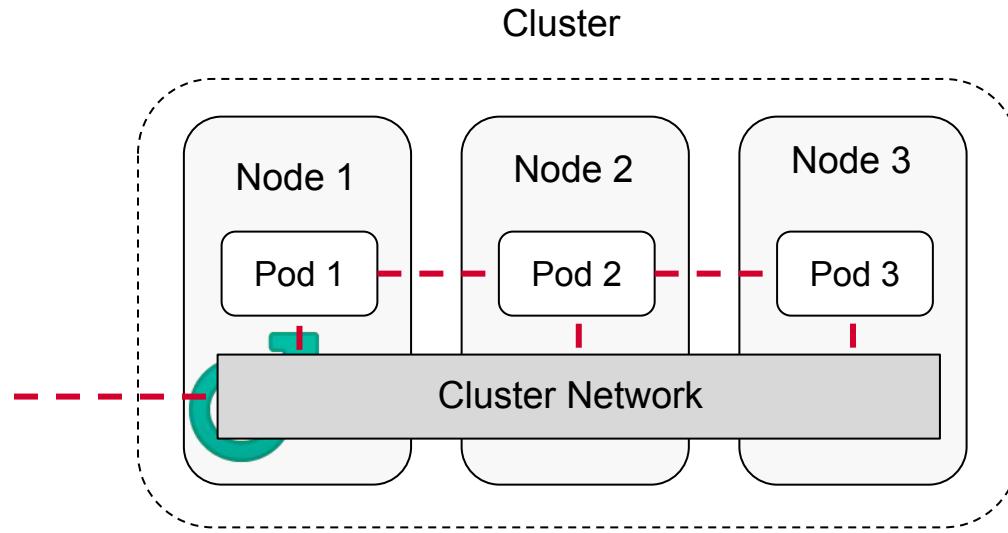
*Follow those guidelines if you are running Kubernetes on AWS*

- Only grant the minimal set of IAM roles required to do the job
- Block access from containers that don't need IAM roles (e.g. via firewall, Network Policies...)



# Scenario: Infiltr. Cluster Network

*The cluster network spans across all nodes*



---

# Minimizing Ext. Network Access

*By default, a Pod can talk to any other Pod in the cluster*

- The cluster network spans across all nodes and enables Pods to communicate with each other
- Access to the cluster networks grants access to all Pods running on all nodes
- Limit network access from the outside



# Linux Security Kernel Modules

*Application that controls what programs can run on host OS*

- AppArmor: Protects the OS from external or internal threats, even zero-day attacks, by enforcing good behavior and preventing flaws from being exploited
- seccomp: Restricts the system calls of an application to the kernel





# AppArmor Installation Check

*Usually exists on Linux system independent of Kubernetes*

```
$ aa-status
apparmor module is loaded.
33 profiles are loaded.
33 profiles are in enforce mode.
...
...
```

Minikube does not support an AppArmor add-on at this time





# AppArmor Profiles

A set of rules that define what a program is allowed/denied to do

```
#include <tunables/global>

profile k8s-deny-write flags=(attach_disconnected)
{
    #include <abstractions/base>

    file,
    # Deny all file writes.
    deny /** w,
}
```





# Enforcing AppArmor Profile

*Store the profile in a file and use apparmor\_parser command*

```
$ sudo apparmor_parser /etc/apparmor.d/k8s-deny-write  
  
$ aa-status  
apparmor module is loaded.  
32 profiles are loaded.  
32 profiles are in enforce mode.  
k8s-deny-write  
...
```

The command uses enforce mode by default. Use -C option for complain mode. **Needs to be enabled on node that runs the Pod.**





# AppArmor Profile in Pod

*Assign annotation, profile needs to be activated on node*

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-apparmor
  annotations:
    container.apparmor.security.beta.kubernetes.io/hello:
      localhost/k8s-allow-write
spec:
  containers:
    - name: hello
      image: busybox:1.28
      command: ["sh", "-c", "echo 'Hello AppArmor!' && sleep 1h"]
```



---

# EXERCISE

Preventing Network  
Calls From a Pod  
With AppArmor

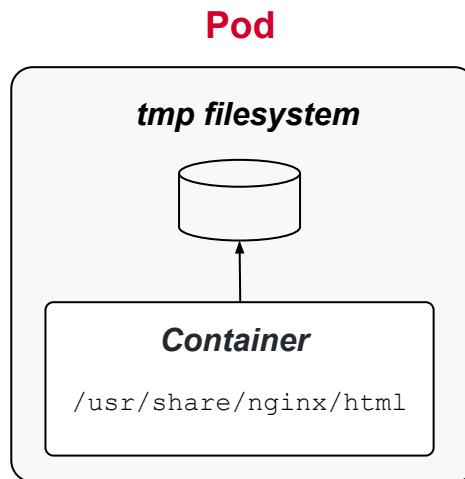


# Minimizing Microservice Vulnerabilities

OS-Level Security Domain, Secrets, Container  
Runtime Sandboxes, mTLS Encryption

# Understanding Security Contexts

*Privilege and access control settings for a Pod or container*



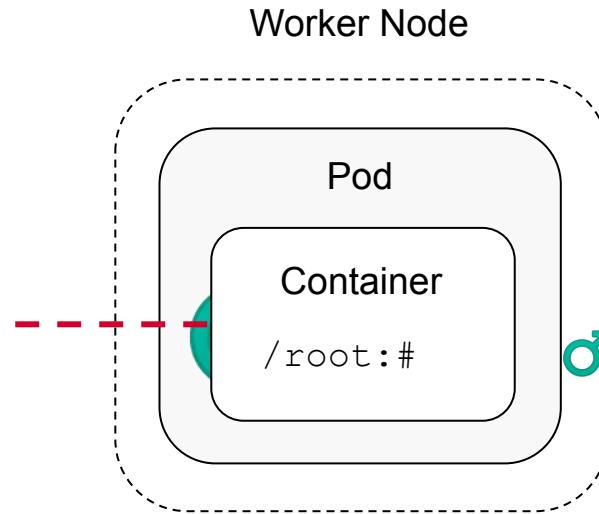
“Create files with a specific Unix group ID”

“Run this container with a specific Unix user ID”



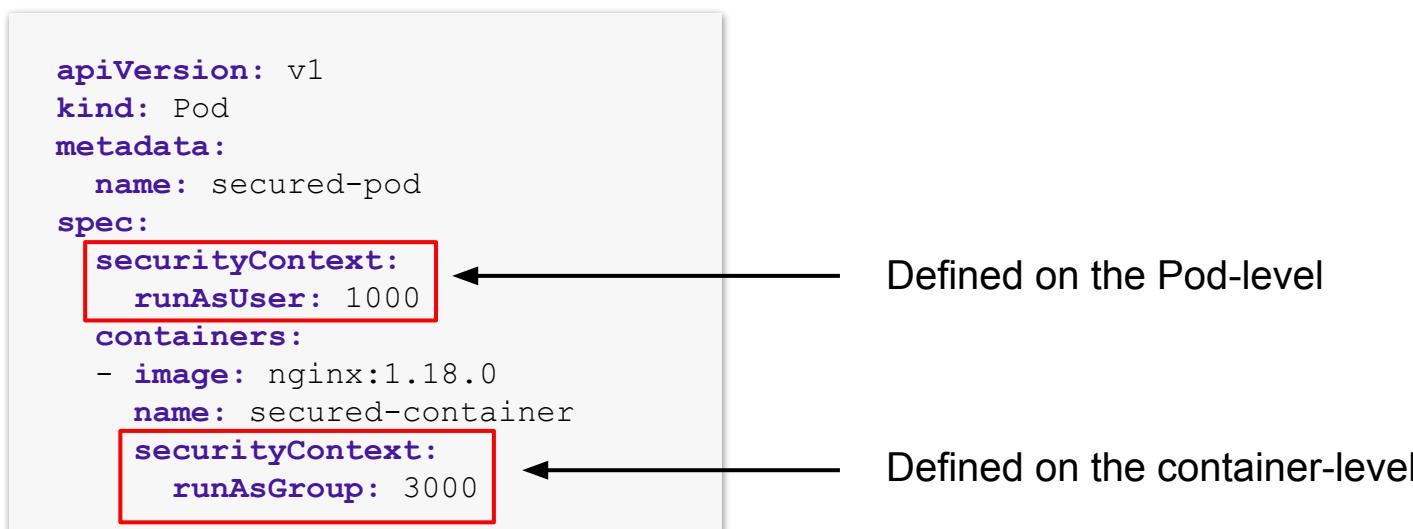
# Scenario: Using root Credentials

*By default, containers run with root privileges*



# Defining a Security Context

*Pod- vs. container-level definition*



# Security Context API

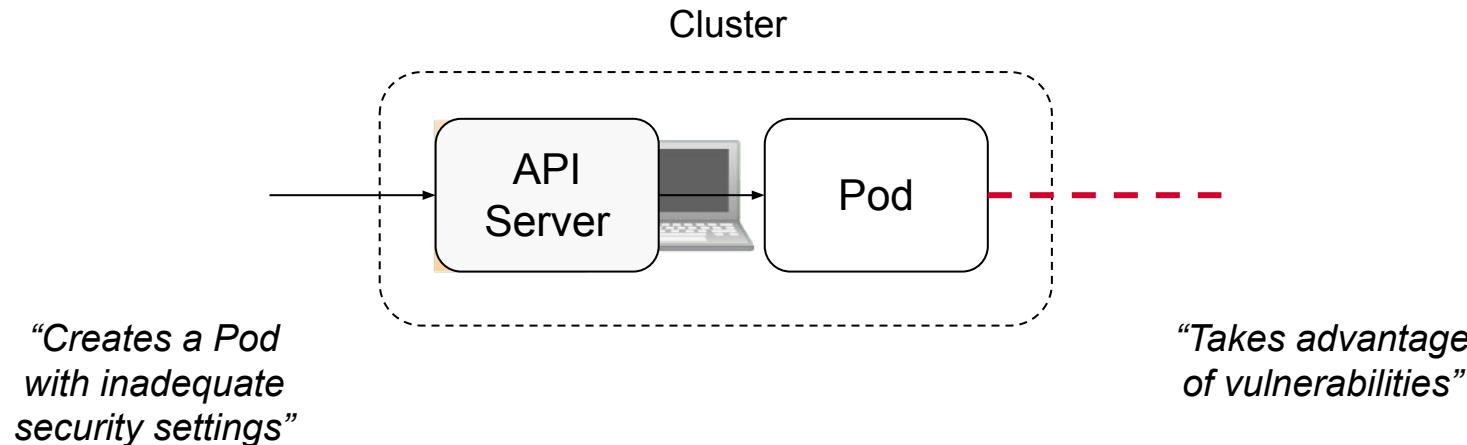
*Only partial overlap for Pod- and container attributes*

| API                                       | Description                                  |
|---|--|
| <a href="#"><u>PodSecurityContext</u></a> | Defines Pod-level security attributes.       |
| <a href="#"><u>SecurityContext</u></a>    | Defines container-level security attributes. |



# Scenario: Governing Pod Security

*Developers may be less familiar with Pod security best practices*



---

# Pod Security Policy Overview

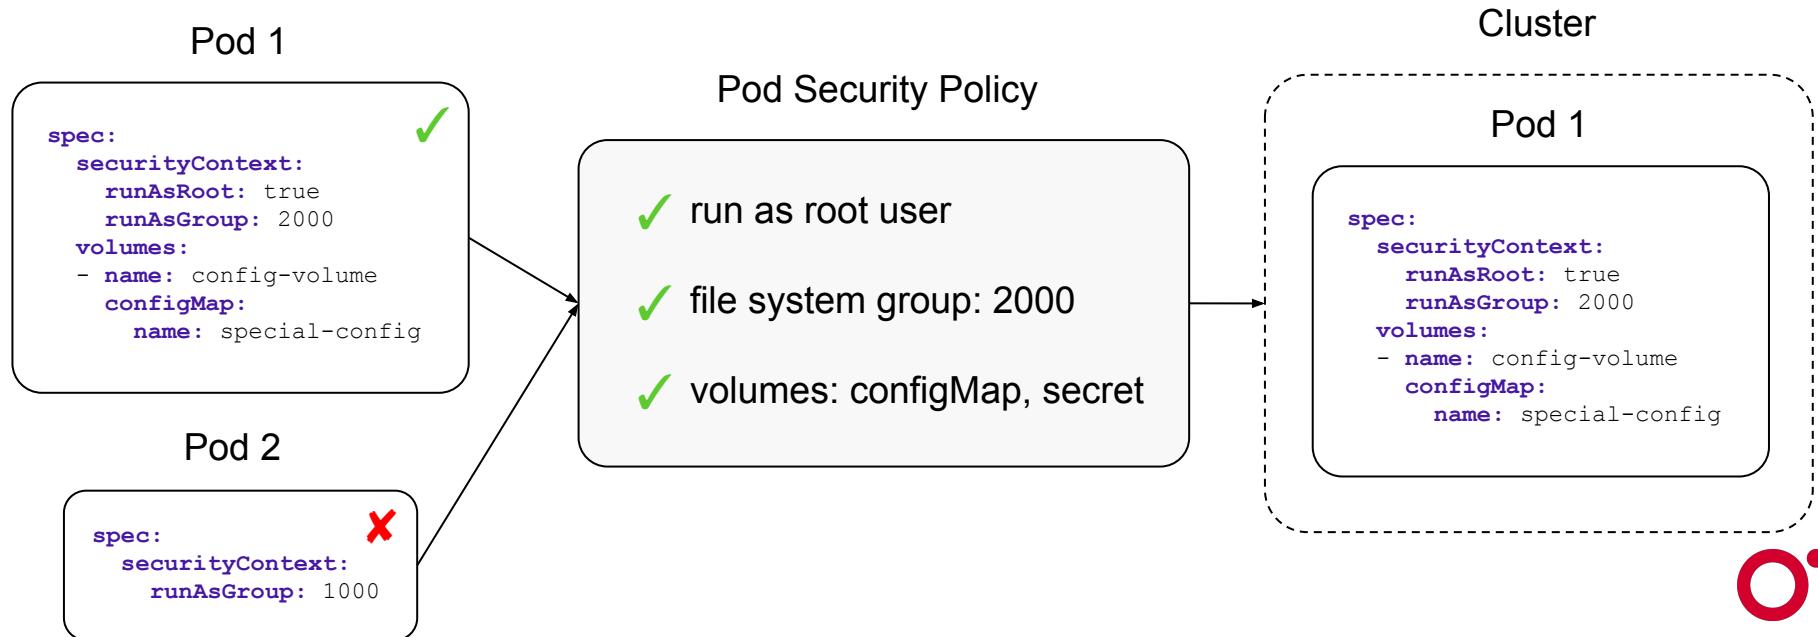
*Administration-level security feature for Pods*

- Controls which security settings should be used for Pods
- Can reject Pods that don't follow the desired security settings
- Examples: Control which types of Volumes are allowed to be mounted, prevent running a container privileged mode or as root user



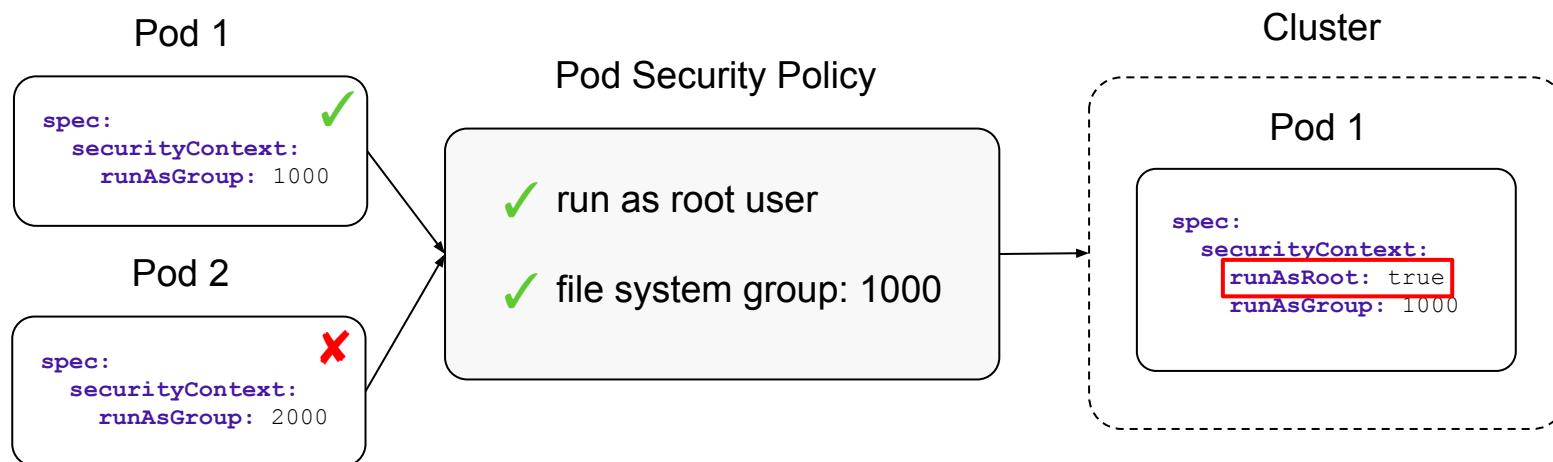
# Pod Security Policy Enforcement

*Requires Pods to set desired security configuration*



# Pod Security Policy Governance

*Automatically assign default security configuration or reject*



# Deprecation Warning

*Still relevant to CKS exam, PSA not in curriculum yet*

## Pod Security Policies

**FEATURE STATE:** Kubernetes v1.21 [deprecated]

**Caution:** PodSecurityPolicy is deprecated as of Kubernetes v1.21, and **will be removed in v1.25**. We recommend migrating to [Pod Security Admission](#), or a 3rd party admission plugin. For a migration guide, see [Migrate from PodSecurityPolicy to the Built-In PodSecurity Admission Controller](#). For more information on the deprecation, see [PodSecurityPolicy Deprecation: Past, Present, and Future](#).



# PSP Admission Controller

*PodSecurityPolicy admission plugin* needs to enabled



```
$ vim /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
      - kube-apiserver
      - ...
      - --enable-admission-plugins=NodeRestriction,PodSecurityPolicy
```

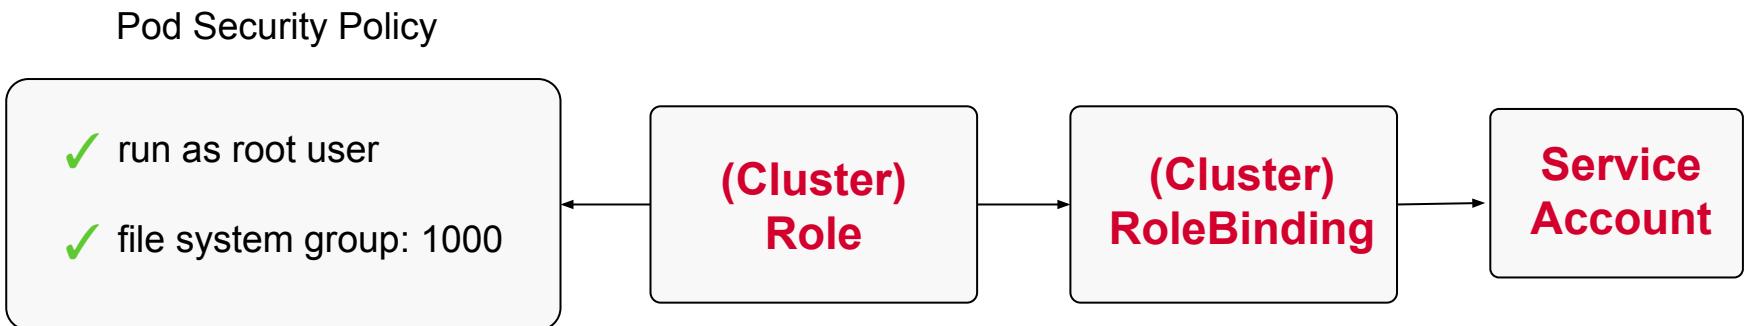


```
$ minikube start
--extra-config=apiserver.enable-admission-plugins,PodSecurityPolicy
--addons=pod-security-policy
```



# Pod Security Policy & RBAC

*To let a policy take effect, you need to tie it in with RBAC*



# Defining a Pod Security Policy

*No imperative command available, extensive configuration*

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: non-privileged
spec:
  privileged: false
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  runAsUser:
    rule: RunAsAny
  fsGroup:
    rule: RunAsAny
  volumes:
  - '*'
```

← List of policies in [reference](#)



# Defining the ClusterRole

*Rule is pretty much set, assign the resource name*

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: psp-clusterrole
rules:
- apiGroups: ['policy']
  resources: ['podsecuritypolicies']
  verbs:     ['use']
  resourceNames:
  - non-privileged
```

← Name of Pod Security Policy



# Defining the (Cluster)RoleBinding

*Apply service account via Group subject, ignore others*

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: psp-clusterrolebinding
roleRef:
  kind: ClusterRole
  name: psp-clusterrole
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: Group
  apiGroup: rbac.authorization.k8s.io
  name: system:serviceaccounts:<authorized namespace>
```

[Recommended service account configuration](#)  
using the Group kind



---

# EXERCISE

Enforcing a Pod  
Security Policy Upon  
Pod Creation





# OPA Gatekeeper

*Customization for Kubernetes objects at creation time*

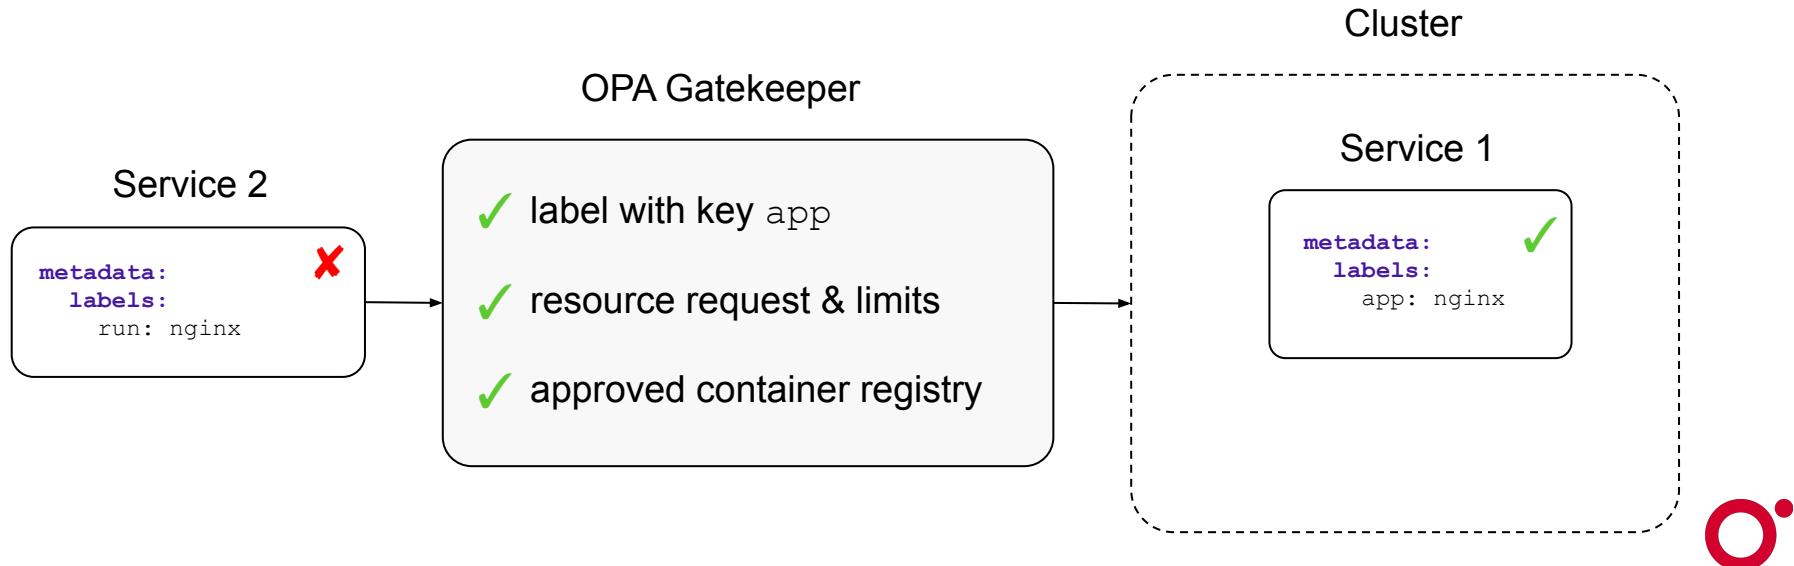
- Framework for defining constraints for any kind of objects, not just Pods
- Highly-flexible policy definition
- Examples: Ensure that certain labels are defined, which image registry are allowed to be used for pulling images





# OPA Enforcement

*Require the definition of a specific label key*





# OPA Key Elements

*Three key elements important to understanding concept*

- Constraint template: Defines the schema for the constraint with the help of Rego logic
- Constraint: Assigns the template to a set of primitives
- Gatekeeper: Components installed in the cluster that enforce the constraints.





# OPA Constraint Template

*No imperative command available, extensive configuration*

```
apiVersion: templates.gatekeeper.sh/v1beta1
kind: ConstraintTemplate
metadata:
  name: k8srequiredlabels
spec:
  crd:
    spec:
      names:
        kind: K8sRequiredLabels
      validation:
        # Schema for the `parameters` field
        openAPIV3Schema:
          properties:
            labels:
              type: array
              items: string
targets:
- target: admission.k8s.gatekeeper.sh
  rego: |
    package k8srequiredlabels

    deny[{"msg": msg, "details": {"missing_labels": missing}}] {
      provided := {label | input.review.object.metadata.labels[label]}
      required := {label | label := input.parameters.labels[_]}
      missing := required - provided
      count(missing) > 0
      msg := sprintf("you must provide labels: %v", [missing])
    }
```

Required parameters and their data types to be passed in

Constraint logic defined with the declarative language Rego





# OPA Constraint

*Maps constraint logic to specific primitives & defines params*

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sRequiredLabels
metadata:
  name: ns-must-have-hr
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Namespace"]
parameters:
  labels: ["hr"]
```

Name of constraint template

Primitives the constraint applies to

Values for parameters





# OPA Gatekeeper Installation

*Deployable with a single command*

**Deploying a release using prebuilt image**

```
$ kubectl apply -f  
https://raw.githubusercontent.com/open-policy-agent/gatekeeper  
/release-3.7/deploy/gatekeeper.yaml
```

**Ensure that Gatekeeper Pods are running**

```
$ kubectl get pods -n gatekeeper-system
```





# OPA Gatekeeper Validation

*Enforcement during object creation*

```
$ kubectl create ns hello
Error from server ([ns-must-have-hr] you must provide labels:
{"gatekeeper"}): admission webhook "validation.gatekeeper.sh"
denied the request: [ns-must-have-hr] you must provide labels:
{"hr"}
```



---

# EXERCISE

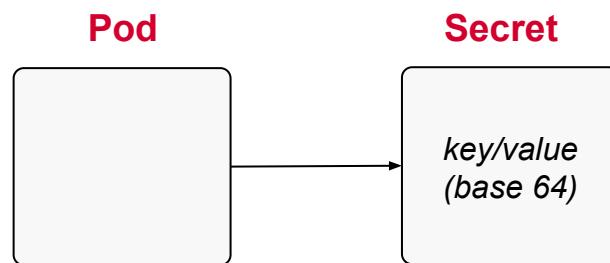
OPA Gatekeeping for  
Pods



---

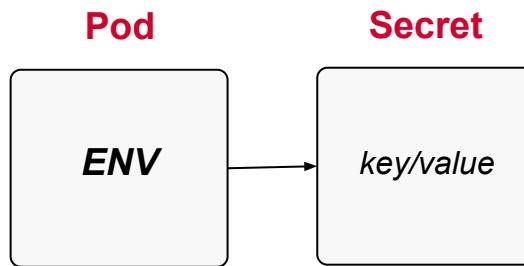
# Centralized Configuration Data

*Stores sensitive data, consumed by Pod*

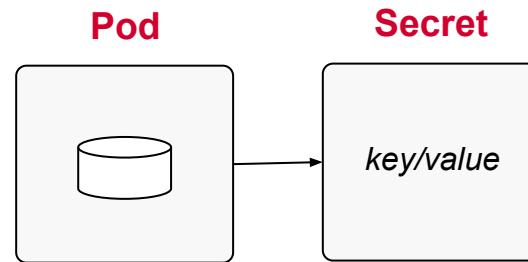


# Mounting a Secret

*Two options for consuming data*



**Injected as environment variables**



**Mounted as volume**



# Secret Options

*Imperative command:* `kubectl create secret`

| Option                       | Description  |
|------------------------------|--|
| <code>generic</code>         | Creates a secret from a file, directory, or literal value. |
| <code>docker-registry</code> | Creates a secret for use with a Docker registry.           |
| <code>tls</code>             | Creates a TLS secret.                                      |



# Creating Secrets (imperative)

*Similar usage to creation of ConfigMap*

```
# Literal values
$ kubectl create secret generic db-creds ←
  --from-literal=pwd=s3cre!

# File containing environment variables
$ kubectl create secret generic db-creds ←
  --from-env-file=secret.env

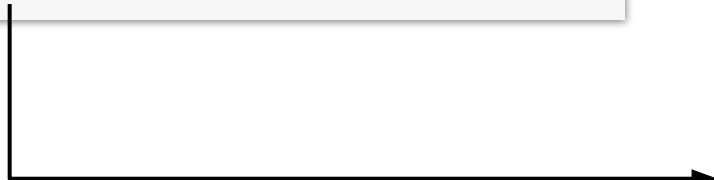
# SSH key file
$ kubectl create secret generic db-creds ←
  --from-file=ssh-privatekey=~/ssh/id_rsa
```



# Creating Secrets (declarative)

*Value has to be base64-encoded manually*

```
$ echo -n 's3cre!' | base64  
czNjcmUh
```



```
apiVersion: v1  
kind: Secret  
metadata:  
  name: mysecret  
type: Opaque  
data:  
  pwd: czNjcmUh
```



# Secret in Pod as Volume

*Value has to be base64-encoded manually*

```
apiVersion: v1
kind: Pod
metadata:
  name: backend
spec:
  containers:
    - name: backend
      image: nginx
      volumeMounts:
        - name: secret-volume
          mountPath: /etc/secret
  volumes:
    - name: secret-volume
      secret:
        secretName: mysecret
```

```
$ kubectl exec -it backend -- /bin/sh
# ls /etc/secret
pwd
# cat /etc/secret/pwd
s3cre!
```



# Specialized Secret Types

*The table only lists some, more in [docs](#)*

| Type                                | Description                          |
|-------------------------------------|--------------------------------------|
| kubernetes.io/basic-auth            | Credentials for basic authentication |
| kubernetes.io/ssh-auth              | Credentials for SSH authentication   |
| kubernetes.io/service-account-token | ServiceAccount token                 |



# Plain-Text Secret Values

*The `stringData` attribute allows for plain-text values*

```
apiVersion: v1
kind: Secret
metadata:
  name: secret-basic-auth
type: kubernetes.io/basic-auth
stringData:
  username: admin
  password: t0p-Secret
```



---

# Detailed Coverage

*The CKAD exam lays the foundation for ConfigMaps/Secrets*

- Certified Kubernetes Application Developer (CKAD) Study Guide
- Exercise: Creating and using a Secret



---

# EXERCISE

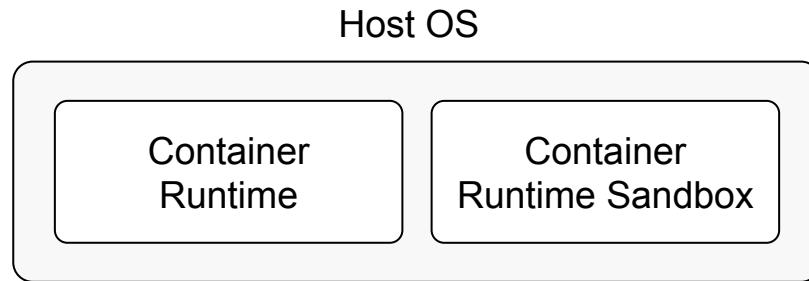
Creating and  
Consuming a Secret



# Container Runtime Sandboxes

*Provides extra layers of process isolation & security*

- Does not execute in standard container runtime (for regular workload)
- The container runtime sandbox exists alongside the container runtime in isolation (for workload we are concerned about)



---

# Container Runtime Sandboxes

*Why would you want to use them?*

- You don't trust the workload as it could execute malicious operations
- Applications that don't need host access
- Provide an extra level of security for workloads executed in a multi-tenant environment (e.g. different customers)



# Implementation Options

*Two solutions are relevant to the exam*



gVisor/runsc: Linux application kernel by Google. Runs within host OS system with a less-lightweight solution than Virtual Machines. This blog post explains the motivation and some implementation details.



Kata containers: Containers run in lightweight Virtual Machines.



---

# Using a gVisor Sandbox

*Three-step process*

- Install the container runtime sandbox implementation e.g. `runsc`
- Configure the container runtime (e.g. Docker Engine, containerd) to interact with the sandbox implementation (see [docs for Docker Engine](#))
- Define a `RuntimeClass` object and reference its name in a Pod



# Runtime Class Definition & Usage

*Defines the container runtime sandbox implementation*

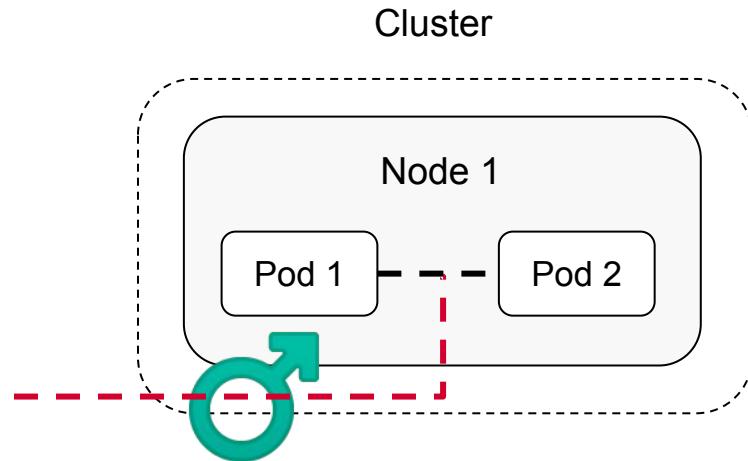
```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc
```

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine
spec:
  runtimeClassName: gvisor
  containers:
    - name: alpine
      image: alpine/curl:3.14
```



# Scenario: Pod-to-Pod Communic.

*Attacker can infiltrate unencrypted Pod to Pod messaging*



---

# Encryption with mTLS

*Like HTTPS but both parties authenticate*

- In the case of Pod-to-Pod communication each Pod has a certificate
- Makes it much harder to impersonate client- or server-side of communication
- Kubernetes API is able to programmatically retrieve a certificate from a Certificate Authority (CA)



# Certificate Process

*More information in [documentation](#)*

- Create a `CertificateSigningRequest` object to request certificate
- Approve or deny the certificate request via `kubectl certificate`
- The attribute `status.certificate` of the `CertificateSigningRequest` object will contain the approved certificate

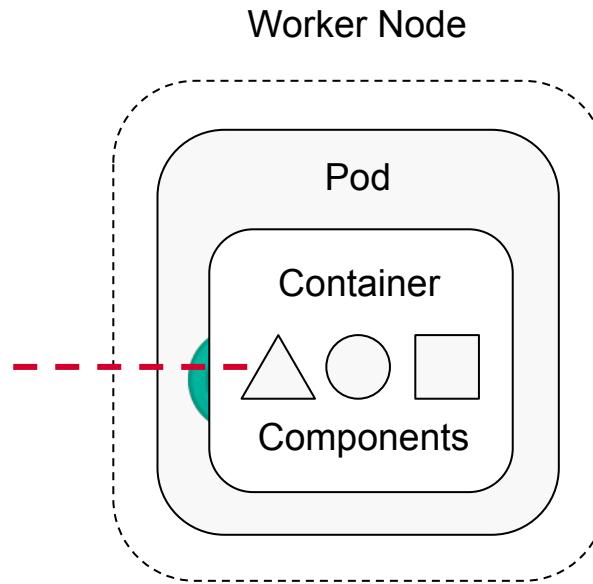


# Supply Chain Security

Base Image Footprint, Image Signing, Validation,  
and Registries, Static Analysis of Workloads,  
Scanning for Vulnerabilities

# Scenario: Exploit Vulnerabilities

*A base image can contain components with known issues*



---

# Minimizing Base Image Footprint

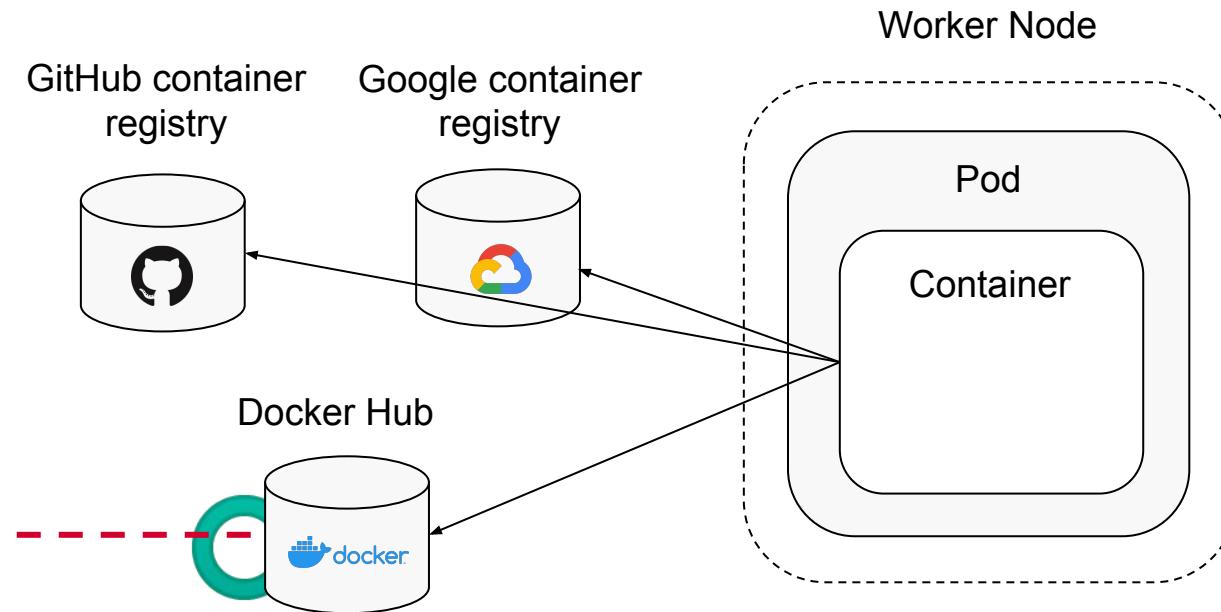
You can find helpful tips in this [blog post](#)

- Choose base images small in size (e.g. alpine uses a minimal OS footprint, distroless don't have a shell)
- Keep your images up to date including their dependencies + configs
- Avoid unnecessary components to be installed (which increases potential security risks) e.g. by using multi-stage builds



# Scenario: Compromised Registry

*Attacker can upload images with malicious instructions*



# Whitelisting Registries

*Only trust an inhouse registry within the network*

- All requests for retrieving container images should be funneled through a trusted registry e.g. Artifactory
- Images can be scanned automatically or upon request
- Enforce the use of a whitelisted registry on Kubernetes e.g. with [OPA Gatekeeper policy](#)



# Validating Images

*Many images have a corresponding hash in registry*

- Hash (e.g. SHA256) can be used to validate the image against
- If validation fails then the image has been manipulated by an attacker
- Append @sha256:<hash> when declaring the image in a Pod



# Image with SHA256

*Definition with an appended hash for validation*

```
apiVersion: v1
kind: Pod
metadata:
  name: alpine
spec:
  containers:
  - name: alpine
    image: alpine:3.15.4@sha256:a777c9c66ba177ccfea2...
```



# Image Validation in Action

*Node fails validation and doesn't download the image*

```
$ kubectl get pods
NAME      READY      STATUS           RESTARTS      AGE
alpine    0/1       ErrImagePull      0            3m12s

$ kubectl describe pod alpine
Events:
  Type      Reason     Age      From          Message
  ----      ----      ----      ----          -----
  Warning   Failed    5s       kubelet        Failed to pull image
"alpine:3.15.4@sha256:a877c9c66ba177ccfea23f2a216ff6721e78a662cd17019488c417135299
cd89": rpc error: code = Unknown desc = Error response from daemon: manifest for
alpine@sha256:a877c9c66ba177ccfea23f2a216ff6721e78a662cd17019488c417135299cd89 not
found: manifest unknown: manifest unknown
```



# Creating a Container Image



*Definition of instructions in a Dockerfile*

- Writing a Dockerfile is “easy” but can introduce issues with sizing, and potential attack vectors
- There are a lot of best practices to adhere to that are not checked automatically



---

# Some Best Practices



*Definition of instructions in a Dockerfile*

- The base image should use a concrete tag other than latest
- Avoid running the container with the `root` user
- Do not use sensitive data anywhere in the Dockerfile
- Use multi-stage builder to avoid having to ship with unnecessary tools



# Creating a Kubernetes Object



*Imperative or declarative approach*

- The declarative approach requires a JSON or YAML manifest file via the `kubectl apply` or `kubectl create` command
- The YAML manifest applies to specific Kubernetes primitive and is very configurable
- Some primitives allow for granting permissions that can be used against you by an attacker





# Some Best Practices

*Restrict permissions, declare resource needs*

- Avoid using privileged mode for Pods if you can avoid it
- Enforce running a container with the `root` user via a security context
- Don't use the `latest` tags for images run in a container
- Define resource requests and limits for containers



---

# Scanning Images

*Detecting vulnerabilities in an automated fashion*

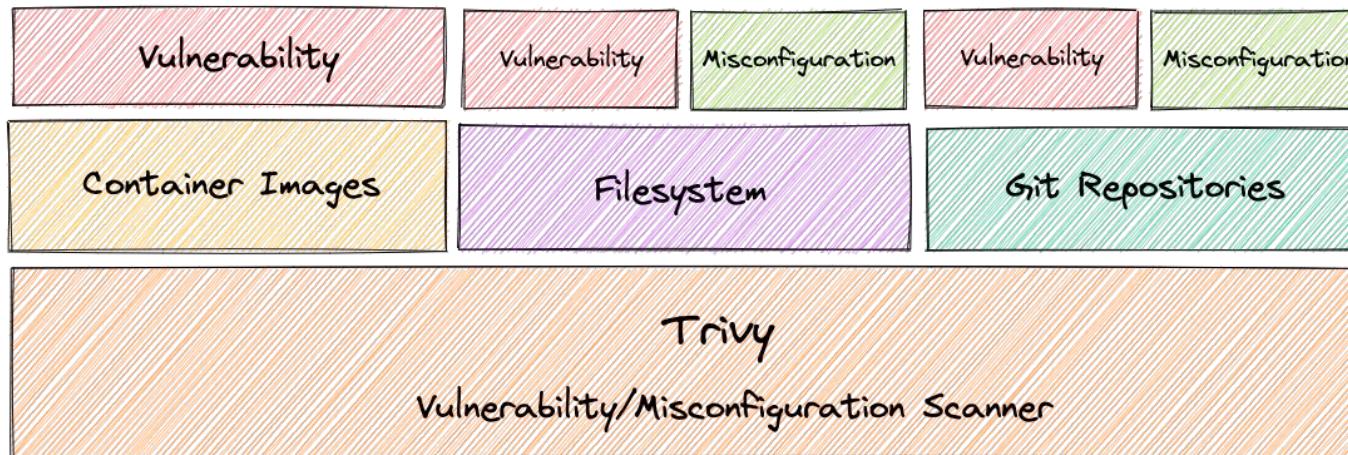
- Vulnerabilities are reported and stored in a central database
- The Common Vulnerability Scoring System (CVSS) assign the degree of severity
- Tools like [Trivy](#) can verify an image with the central database of vulnerabilities





# Using Trivy

*CLI tool for scanning images, Git repos, config files, FS*





# Installing & Using Trivy

*Easy to install and use with different modes*

- Installation: apt-get install, yum install **and** brew install
- Scanning an image: \$ trivy image nginx:1.18.0
- Scanning a config file: \$ trivy config pod.yaml
- Scanning a repo: \$ trivy repo  
<https://github.com/bmuschko/cks-crash-course>





# Trivy Vulnerability Report

*Lists vulnerabilities per library used in an image*

```
2022-04-20T11:24:24.060-0600      INFO      Table result includes only package filenames. Use '--format json' option to get the full path to the package file.

Node.js (node-pkg)
=====
Total: 35 (UNKNOWN: 0, LOW: 7, MEDIUM: 12, HIGH: 14, CRITICAL: 2)

+-----+-----+-----+-----+-----+
| LIBRARY | VULNERABILITY ID | SEVERITY | INSTALLED VERSION | FIXED VERSION | TITLE
+-----+-----+-----+-----+-----+
| ajv (package.json) | CVE-2020-15366 | MEDIUM | 5.5.2 | 6.12.3 | nodejs-ajv: prototype pollution
|                   |                   |          |          |          | via crafted JSON schema
|                   |                   |          |          |          | in ajv.validate function
|                   |                   |          |          |          | -->avd.aquasec.com/nvd/cve-2020-15366
+-----+-----+-----+-----+-----+
```



---

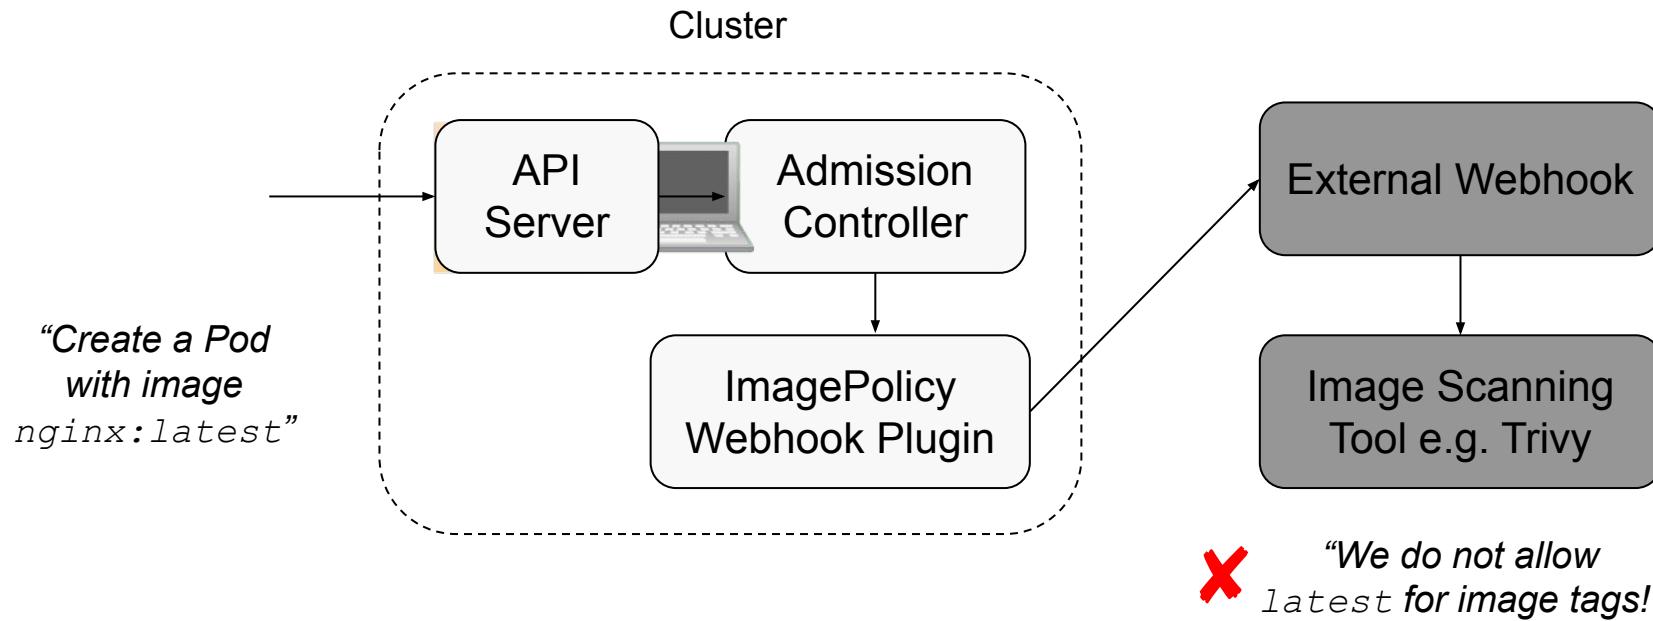
# EXERCISE

Image Scanning with  
Trivy



# Scenario: Intercepting API Calls

*Image approval via admission controller and image policy webhook*



# Admission Controllers

*Intercept calls via the Kubernetes API*

- Decides if a request should be approved/denied or mutated before any changes take place to enforce aspects like security, governance, or configuration management
- Evaluation is based on information sent e.g. image name
- Example: Disallow containers from running with a non-root user when creation of a Pod is requested



---

# ImagePolicyWebhook Plugin

*Make admission decisions via backend*

- One of the [plugins](#) for an admission controller
- Approve/deny usage of image when creating workloads e.g. scan image
- Example: Deny an image that contains known critical vulnerabilities



# Admission Configuration

*Deny Pod creation via ImagePolicyWebhook plugin*

```
apiVersion: apiserver.config.k8s.io/v1
kind: AdmissionConfiguration
plugins:
- name: ImagePolicyWebhook
  configuration:
    imagePolicy:
      kubeConfigFile: /etc/kubernetes/admission/kubeconf
      allowTTL: 50
      denyTTL: 50
      retryBackoff: 500
      defaultAllow: false
```

Deny all Pods if  
backend is not  
available



# Enabling ImagePolicyWebhook

*Point to admission config file and enable plugin*



```
$ vim /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
      - kube-apiserver
      - ...
      - --admission-control-config-file=/etc/kubernetes/admission/admission-config.yaml
      - --enable-admission-plugins=NodeRestriction,ImagePolicyWebhook
```



```
$ minikube start
--extra-config=apiserver.admission-control-config-file=...
--extra-config=apiserver.enable-admission-plugins=ImagePolicyWebhook
```



# The Kubeconfig File

*Sets up communication with backend via TLS*

```
clusters:  
- name: name-of-remote-imagepolicy-service  
  cluster:  
    certificate-authority: /path/to/ca.pem  
    server: https://images.example.com/policy  
  
users:  
- name: name-of-api-server  
  user:  
    client-certificate: /path/to/cert.pem  
    client-key: /path/to/key.pem
```

HTTPS URL of  
remote service to  
query

Cert and key for the  
webhook admission  
controller to use



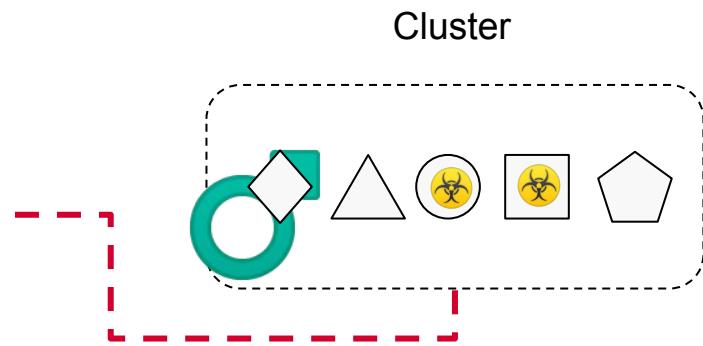
# Monitoring, Logging, and Runtime Security

Behavioral Analysis, Detecting Threats, Immutable Containers, Audit Logs

---

# Scenario: Infiltrating Components

*Detecting if and how cluster components are attacked*



---

# Behavioral Analytics

*Process of observing abnormal & malicious events*

- Answers questions like “Are we currently under attack?” or “Was something changed after an attack?”
- Can be performed in a manual and automated fashion
- Tools like [Falco](#) continuously detect unexpected behavior, configuration changes, intrusions, and data theft in real time





# Installing & Using Falco



*Only installable and executable on Linux systems*

- Modes of operation
  - CLI tool
  - Background process
  - Web server

```
$ falco -r custom-rules.yaml -M 60
```

Length of  
execution time



Custom rules file



# Falco Rules



*Define the conditions that trigger an alert*

```
- rule: shell_in_container
  desc: notice shell activity within a container
  condition: evt.type = execve and evt.dir=< and container.id != host and proc.name = bash
  output: shell in a container (user=%user.name
            container_id=%container.id
            container_name=%container.name
            shell=%proc.name parent=%proc.pname
            cmdline=%proc.cmdline)
  priority: WARNING
```

*“Alert whenever a user opens a shell to a container”*



# Falco 101 Video Course



*Free with online registration at [Sysdig webpage](#)*

A screenshot of a web browser displaying the 'Learn with Sysdig' website at learn.sysdig.com. The main header features the Sysdig logo and the text 'Learn with Sysdig'. On the right, there's a 'Sign In' button. The central content area is titled 'Falco 101' with the subtitle 'All you need to learn to get started with Falco'. It includes a 'Register Now' button, a 'Sign In' link for registered users, and social sharing links for Facebook and Twitter. Below this, there's a section for 'ALSO AVAILABLE AS PART OF:' with a small thumbnail image. The main content area is divided into two columns: 'About this course' on the left and 'Curriculum' on the right.

**About this course**

This course introduces Falco, including its rules syntax, alerting, configuration, and more.

**What is Runtime Security?**

**How can I use the default rule set for my particular use case?**

**Can I define custom rules for my production environment?**

If some of these questions resonate with you, this is your course. The first three lessons are composed of a short video introducing the course and main concepts. The rest of the content is hands-on oriented, so you'll have the opportunity to get some actual practice with the tool in real scenarios.

**Curriculum (5 hr 41 min)**

- 101.01 - Introduction (0 hr 1 min)
- 101.02 - Introduction to Runtime Security (0 hr 3 min)
- 101.03 - What is Falco? (0 hr 2 min)
- 101.04 - Falco Installation and basic settings (0 hr 30 min)
- 101.05 - Falco Event Sources (0 hr 0 min)
- 101.06 - Falco Rules Basics (1 hr 0 min)
- 101.07 - Falco Rules Deep Dive (0 hr 45 min)
- 101.08 - Fine Tuning Falco Rules (Exceptions) (0 hr 45 min)
- 101.09 - Falco Alerts (0 hr 35 min)
- 101.10 - Falco Response Engine (0 hr 30 min)

*Check what you learned!*



---

# EXERCISE

Intrusion Detection  
with Falco



---

# Container Immutability

*Cannot be changed during its lifespan*

- Container's file system cannot be modified (no updates, no patches, no configuration changes)
- Updates can only be made via new container image
- Container configuration needs to be externalized e.g. environment variables, ConfigMaps, Secrets and can be used to control runtime behavior



# Avoid Privileged Containers

*No access to all the devices of the host machine*

- Running container with root access opens door for attackers to gain full access to server
- Control host access in the Pod's security context configuration
- Use PodSecurityPolicies to prevent privileged access



---

# Immutability Best Practices

*No access to all the devices of the host machine*

- Avoid running container with root user via `securityContext.runAsUser` with value `root` or `0`
- Avoid privileged mode via `securityContext.privileged` or `securityContext.allowPrivilegedEscalation` with value `true`
- Avoid access to host resources via `hostNetwork` with value `true`



# Ensuring Container Immutability

*Prevent writing to container file system*

```
apiVersion: v1
kind: Pod
metadata:
  name: immutable
spec:
  securityContext:
    readOnlyRootFilesystem: true
  containers:
    - ...
```

*Mount a Volume if you need to write data from the application*



---

# EXERCISE

Verifying Container  
Immutability



# Audit Logging

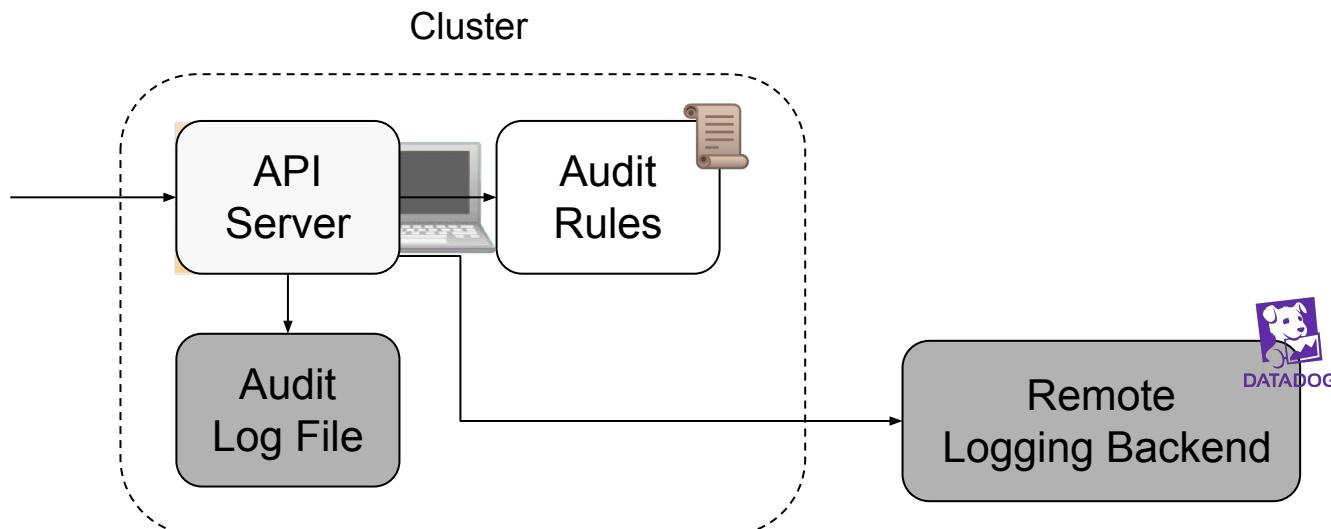
*Chronological records of actions in the cluster*

- Records requests to Kubernetes API for internal and external components (e.g. a user creates a Pod with `kubectl`, or a change to CoreDNS)
- Structured in JSON format
- Helpful for real-time view on cluster events (e.g. threat detection) or why or what happened in the past (e.g. for debugging issues)



# Audit Backends

*Write to file system or external HTTP API*



# Audit Policy Configuration File

*Rules about events and their data to be recorded*

```
apiVersion: audit.k8s.io/v1
kind: Policy
rules:
- level: RequestResponse
  resources:
  - group: ""
    resources: [ "pods" ]
```

Audit level, in this case event metadata,  
request and response bodies

Kubernetes resource types to match  
including a potential API group



# Audit Levels

| Level           | Runtime Behavior   |
|-----------------|--|
| None            | Omit log events  |
| Metadata        | Log request metadata (requesting user, timestamp, resource, verb, etc.) but not request or response body |
| Request         | Log event metadata and request body but not response body  |
| RequestResponse | Log event metadata, request and response bodies  |



# File System Audit Logging

*Point to audit log config file and output log file*



```
$ vim /etc/kubernetes/manifests/kube-apiserver.yaml
spec:
  containers:
    - command:
      - kube-apiserver
      - ...
      - --audit-policy-file=/etc/kubernetes/audit-policy.yaml
      - --audit-log-path=/var/log/kubernetes/audit/audit.log
```



```
$ minikube start
--extra-config=apiserver.audit-policy-file=/etc/ssl/certs/audit-policy.yaml
--extra-config=apiserver.audit-log-path=
```

Workaround: Config  
dir/file needs to exist in  
~/.minikube/files/



# Persisting Audit Log Records

*Requires mounting the config file + produced audit log*

## Mount Volumes

```
volumeMounts:  
- mountPath: /etc/kubernetes/audit-policy.yaml  
  name: audit  
  readOnly: true  
- mountPath: /var/log/kubernetes/audit/  
  name: audit-log  
  readOnly: false
```

## Configure host paths

```
volumes:  
- name: audit  
  hostPath:  
    path: /etc/kubernetes/audit-policy.yaml  
    type: File  
- name: audit-log  
  hostPath:  
    path: /var/log/kubernetes/audit/  
    type: DirectoryOrCreate
```



---

# EXERCISE

Configuring and  
Using Audit Logging



# Summary & Wrap Up

Last words of advice...

---

# Gaining confidence

- Run through practice exams as often as you can
- Read through online documentation start to end
- Know your tools (especially vim, bash, YAML)
- Pick time you are most comfortable, get enough sleep
- Take your first attempt easy but give it your best





Thank you

